Team members: 19C010 - Aswath Swasun P

19C092 - Shane Rex S

19C105 - Subash A

19C117 - Vijay G

Gmeet recording: https://drive.google.com/file/d/1BeX-V97SIXiSoaygFVdxLkux4vELooR1/view

Github link: https://github.com/subash2121/cloud_assignment_3

Cloud Assignment 3

a.

Virtualisation software used: Virtual Box

Virtualisation OS used: Ubuntu

(i) Inter Virtual-machine communication on a single host

Procedure:

1.Create a new Virtual Machine in the Virtual Box

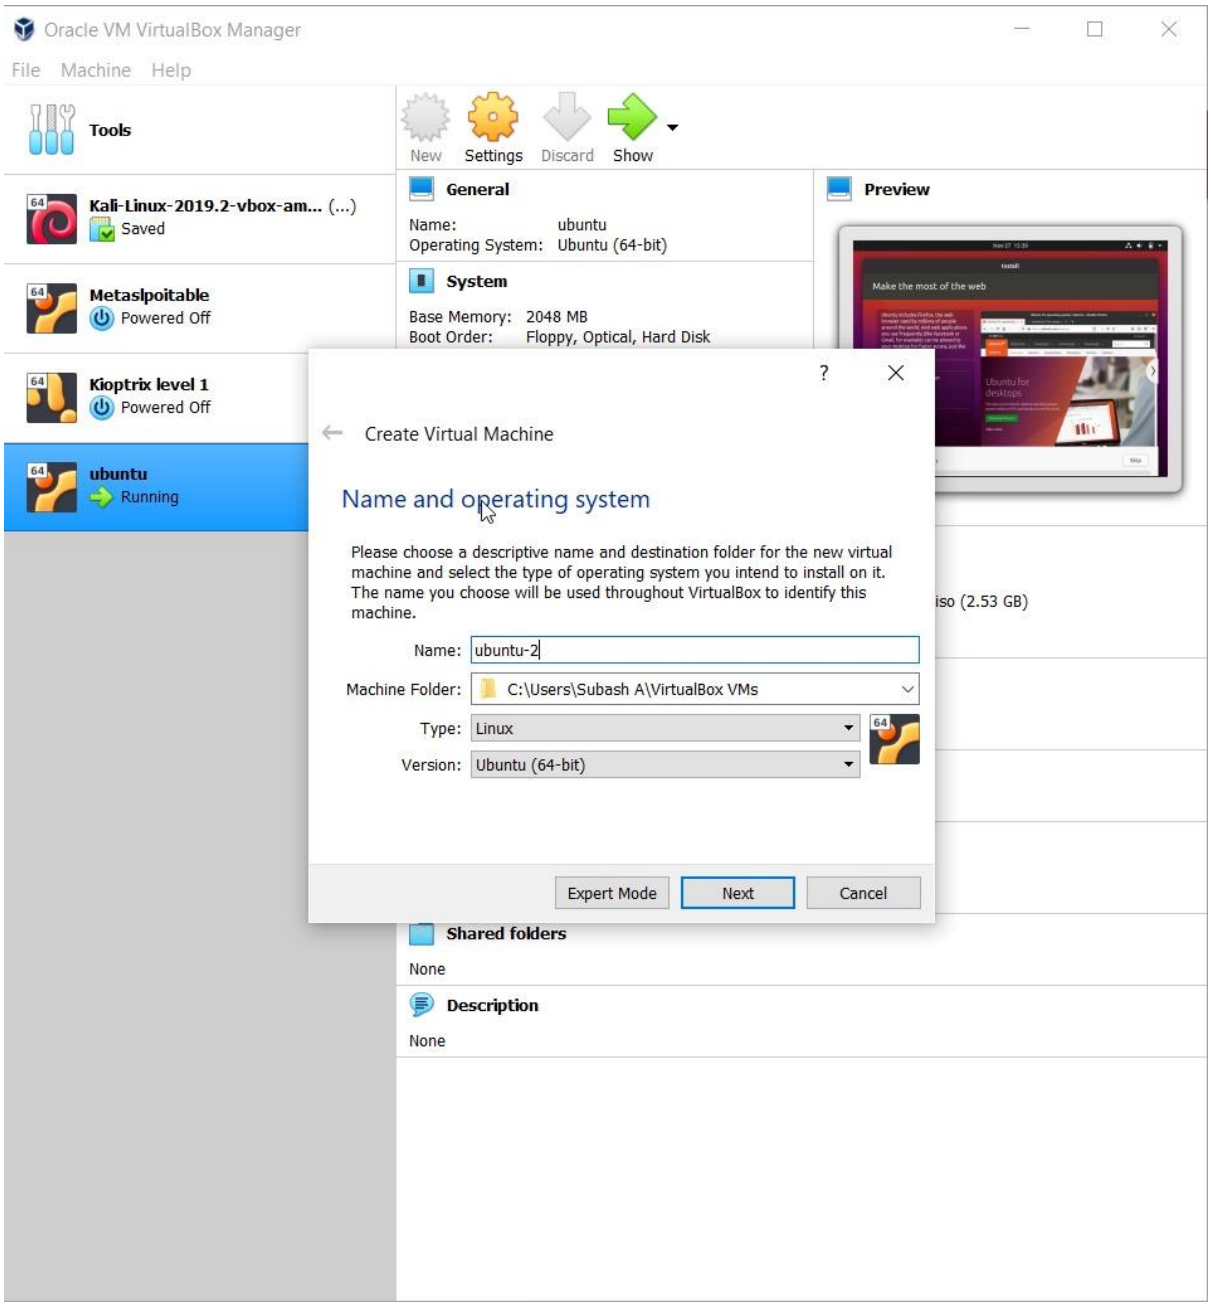

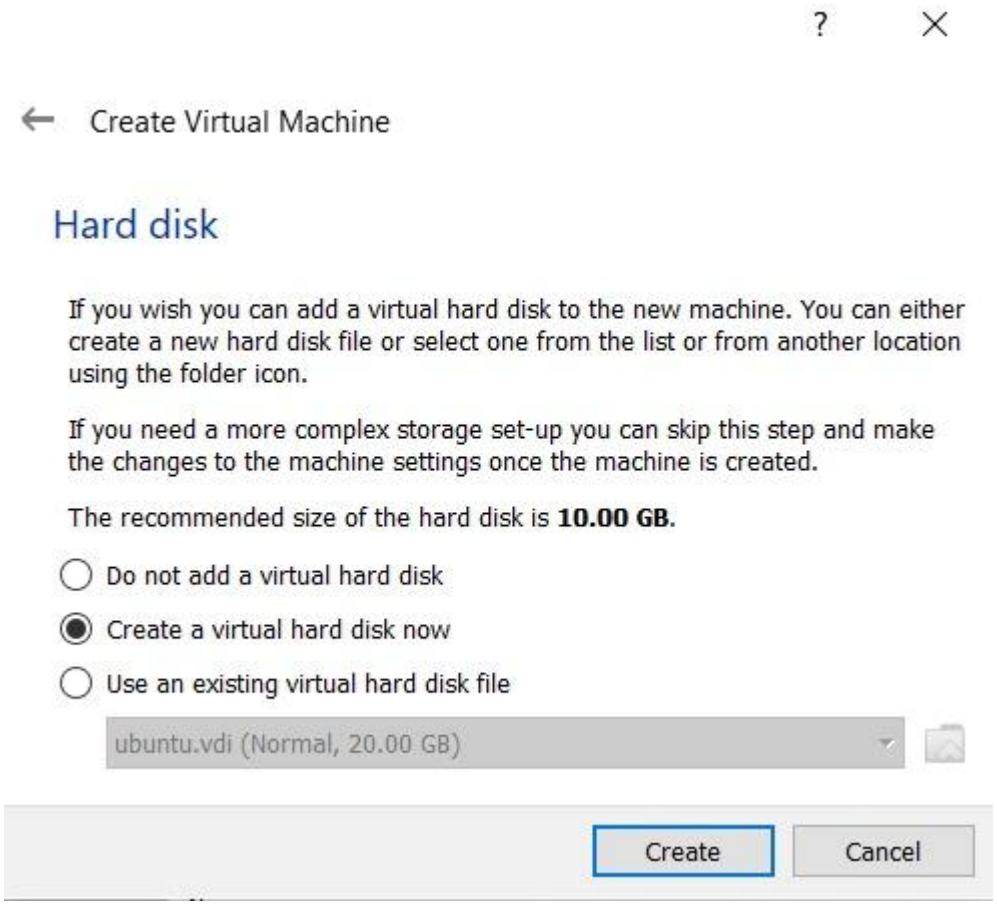
Create Virtual Machine
Hard disk
If you wish you can add a virtual hard disk to the new machine. You can either create a new hard disk file or select one from the list or from another location using the folder icon.
If you need a more complex storage set-up you can skip this step and make the changes to the machine settings once the machine is created.
The recommended size of the hard disk is 10.00 GB.
Do not add a virtual hard disk
Create a virtual hard disk now
Use an existing virtual hard disk file
ubuntu.vdi (Normal, 20.00 GB)
Create
Cancel

2.Allocate the Ram space for the Virtual Machine.

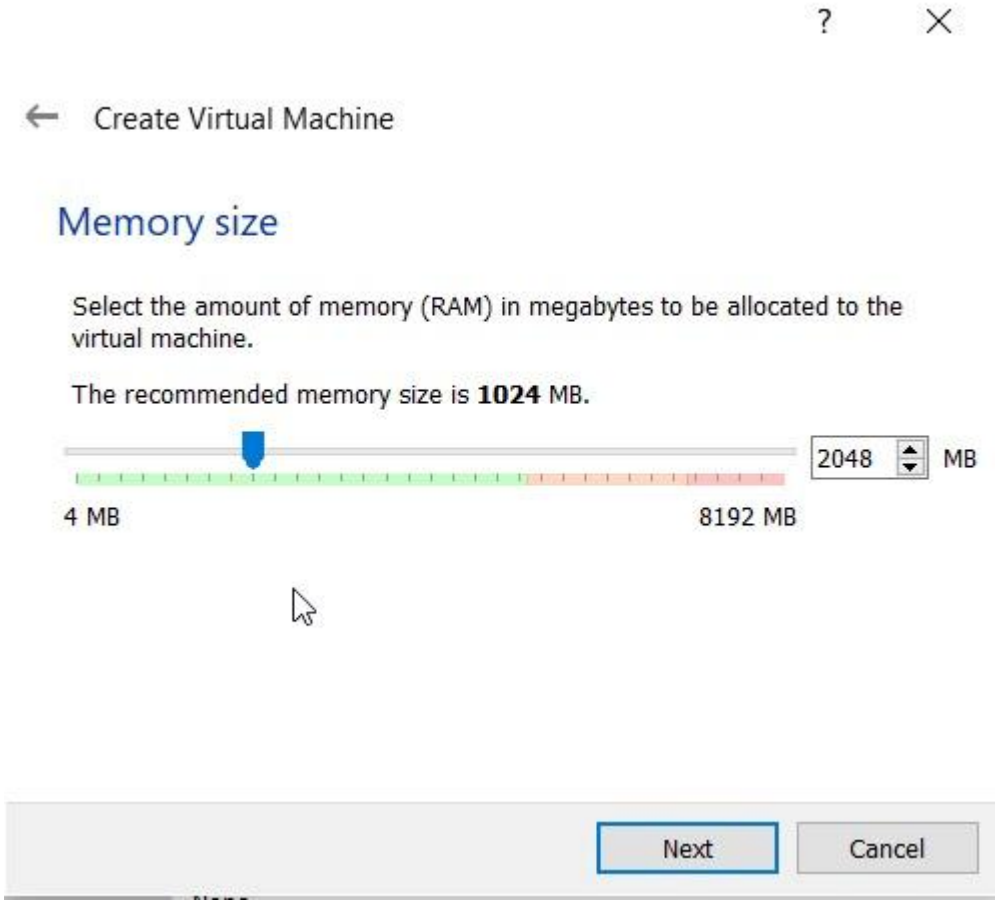

3.Select the Folder in which the Virtual Machine is to be created and the space allocated to the Virtual machine

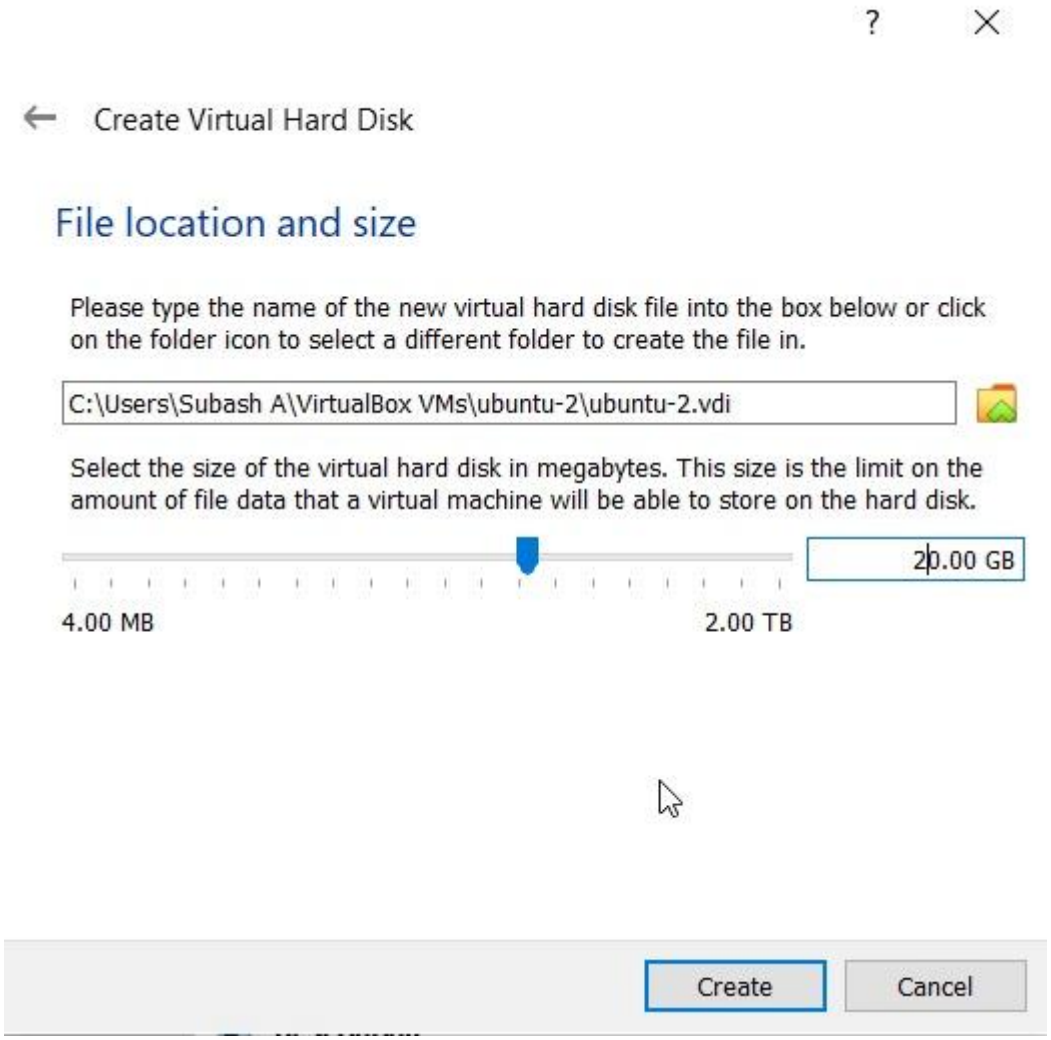

4.The Virtual Machine is now created and now we load the OS ISO file.

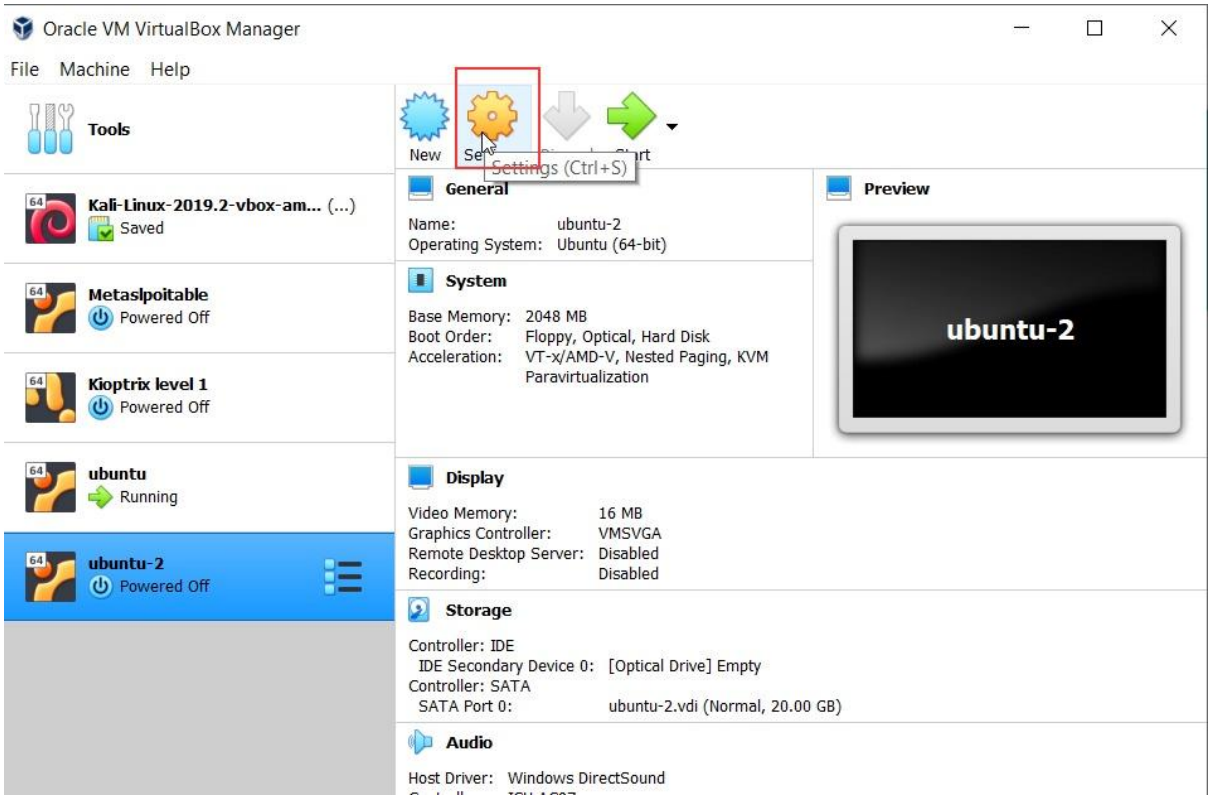

5.Navigate and Select the Ubuntu ISO file.

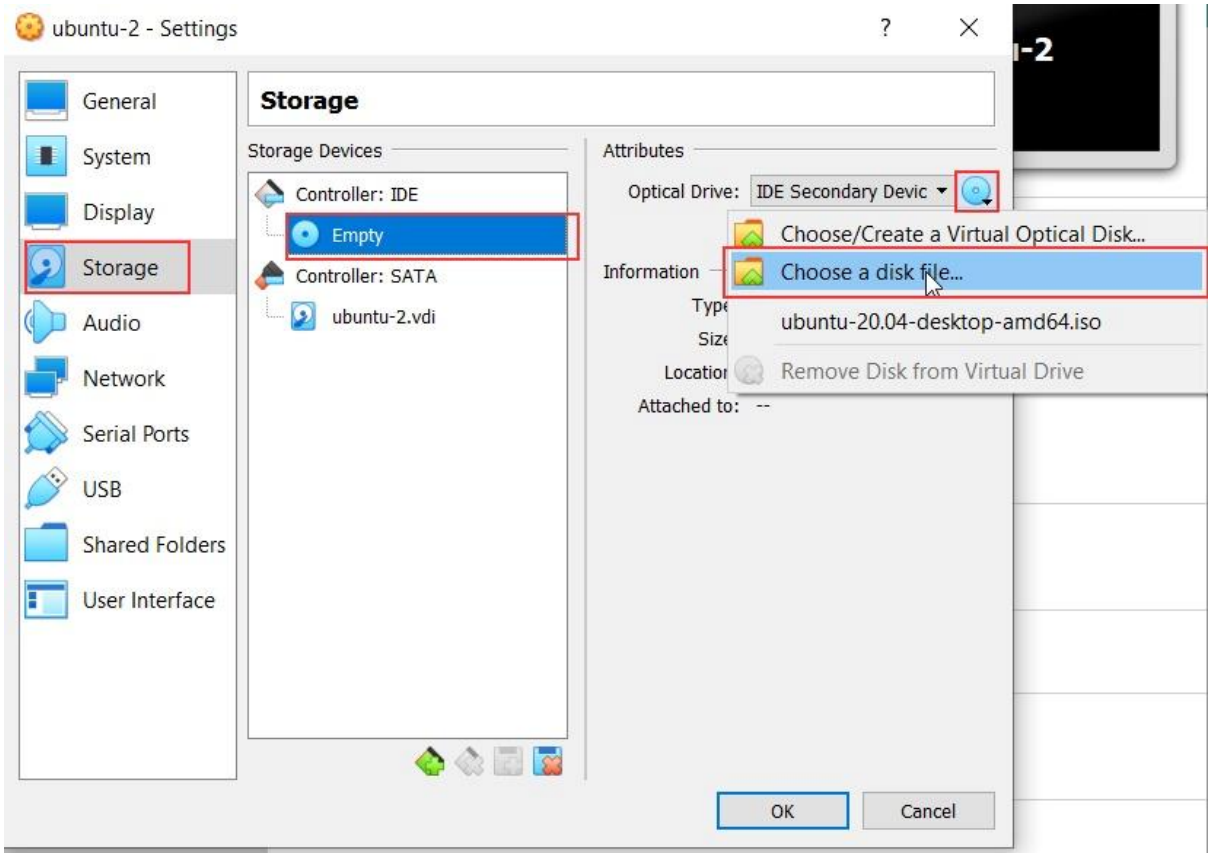

6.Wait for the machine to boot.

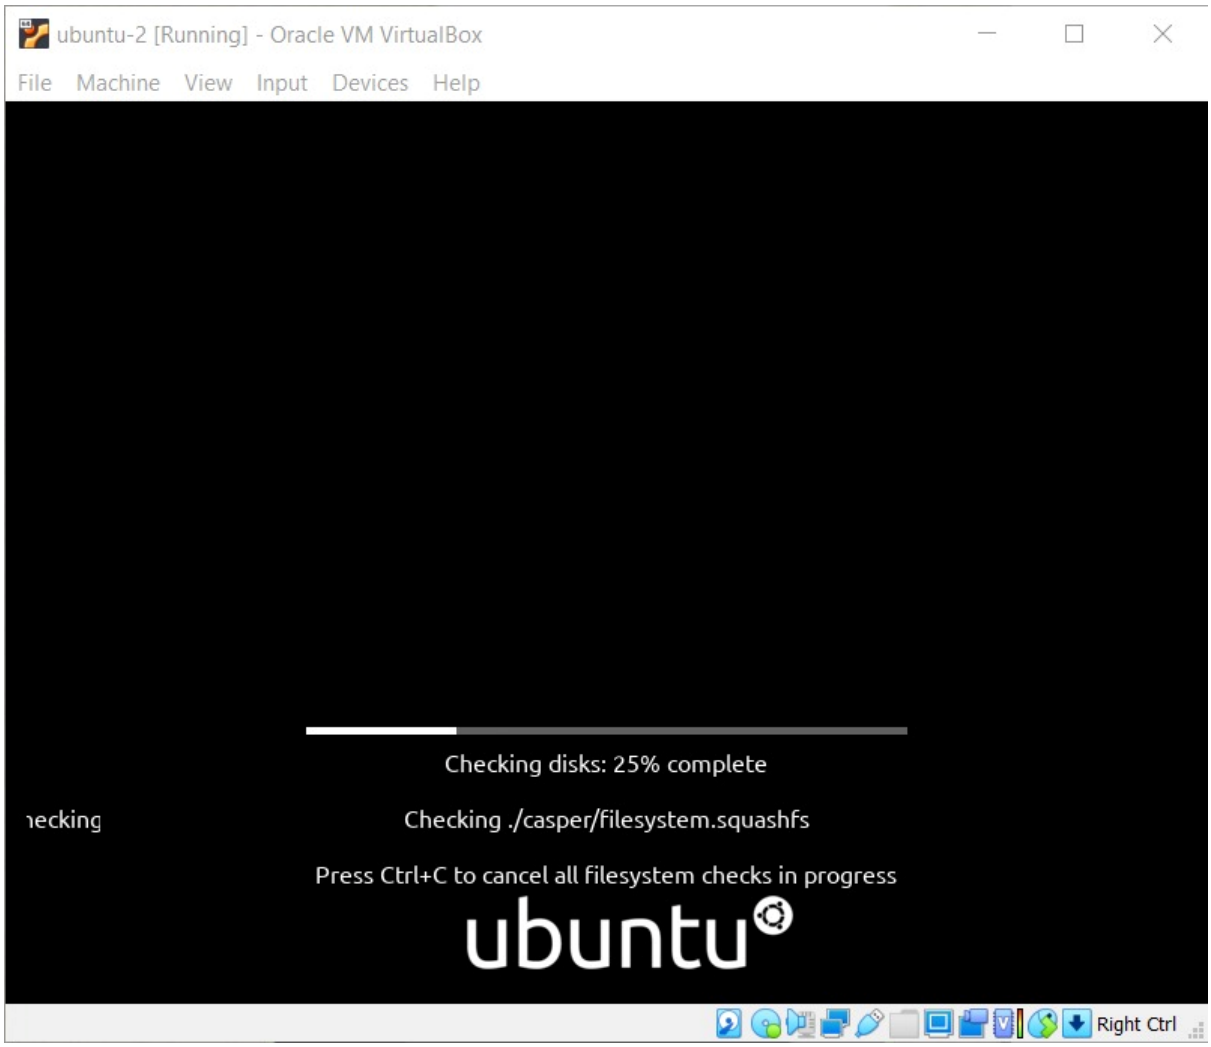

7.Select the way in which you want to run the OS

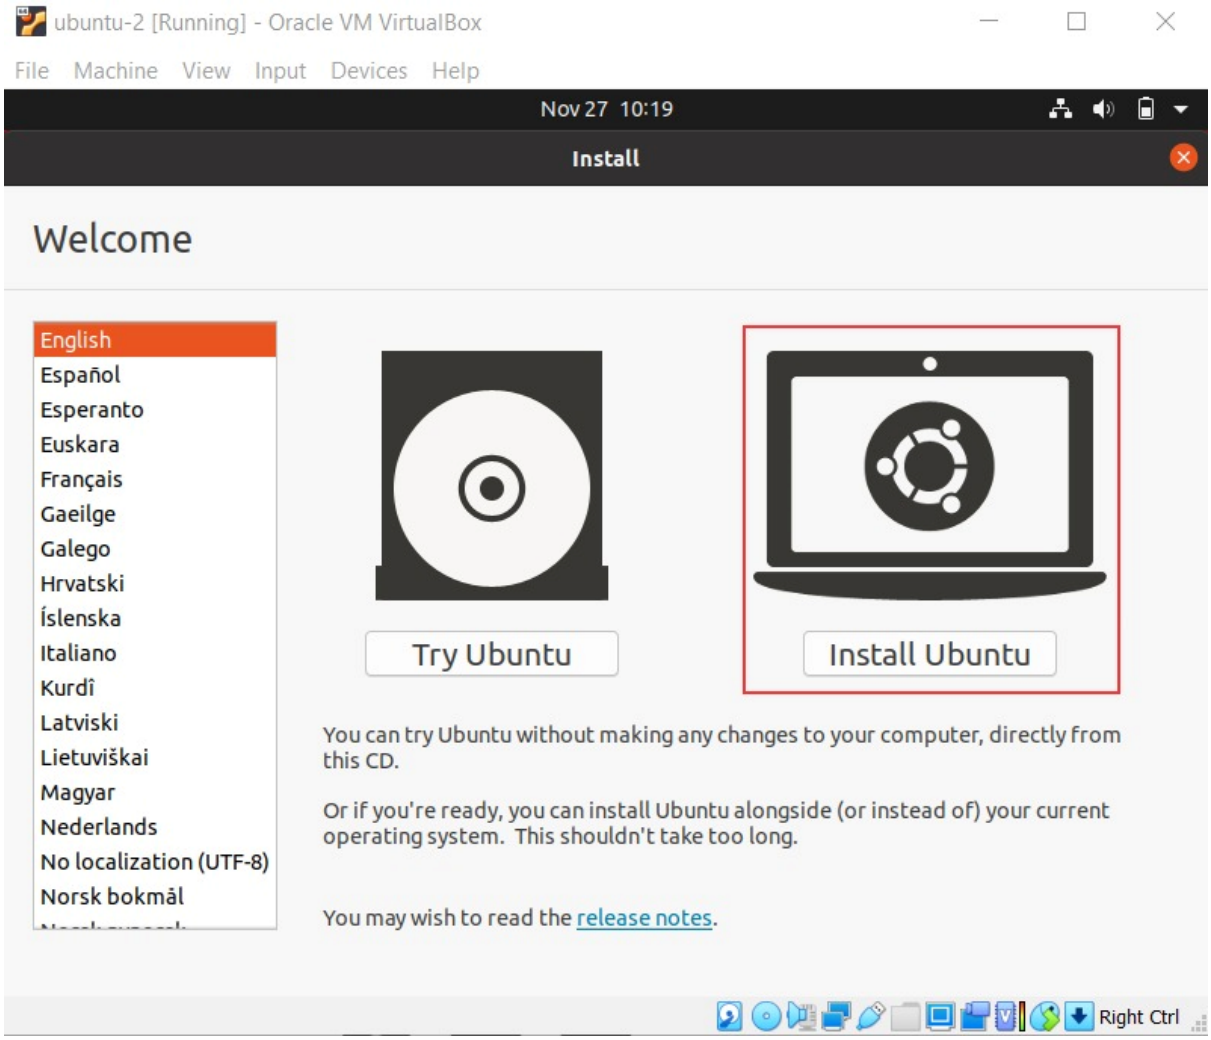

8.Go on with the instructions

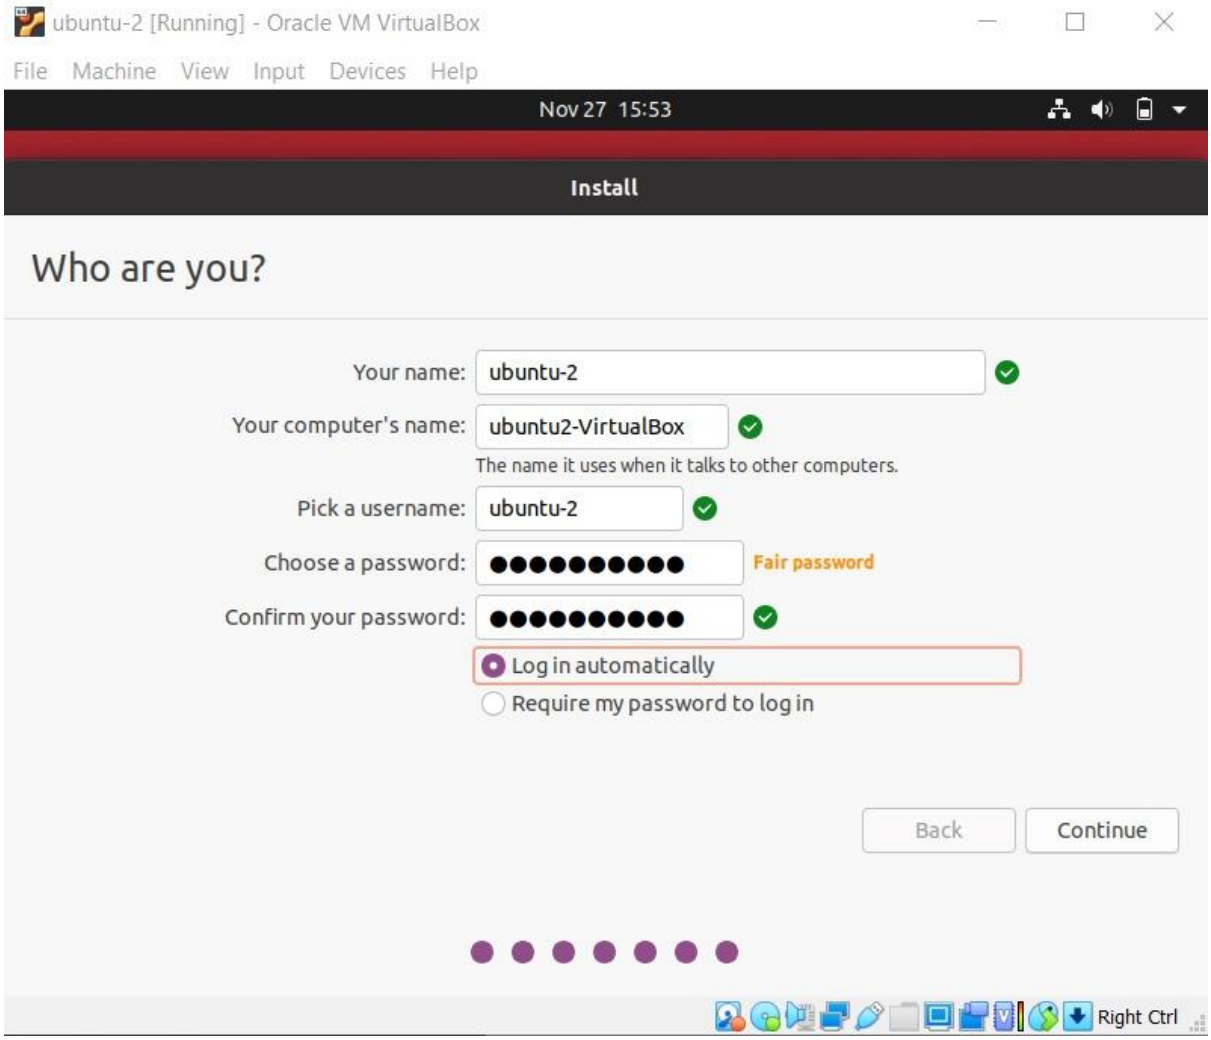

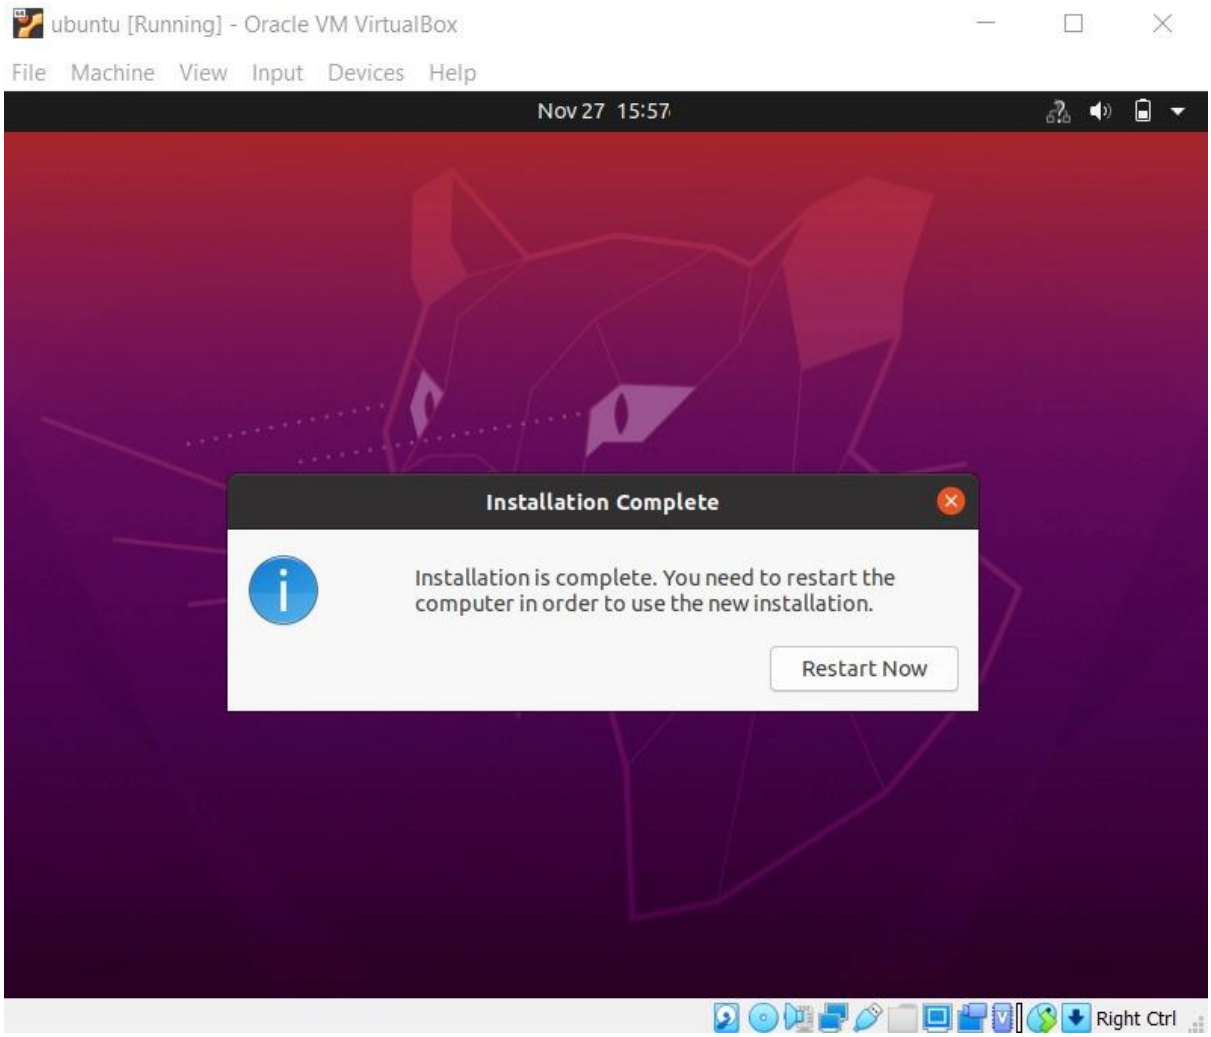

Perform the steps to create a virtual machine inorder to create another machine to perform the Inter Virtual Machine communication.

9.Configure the network inside the vitual box by creating a custom NAT network.

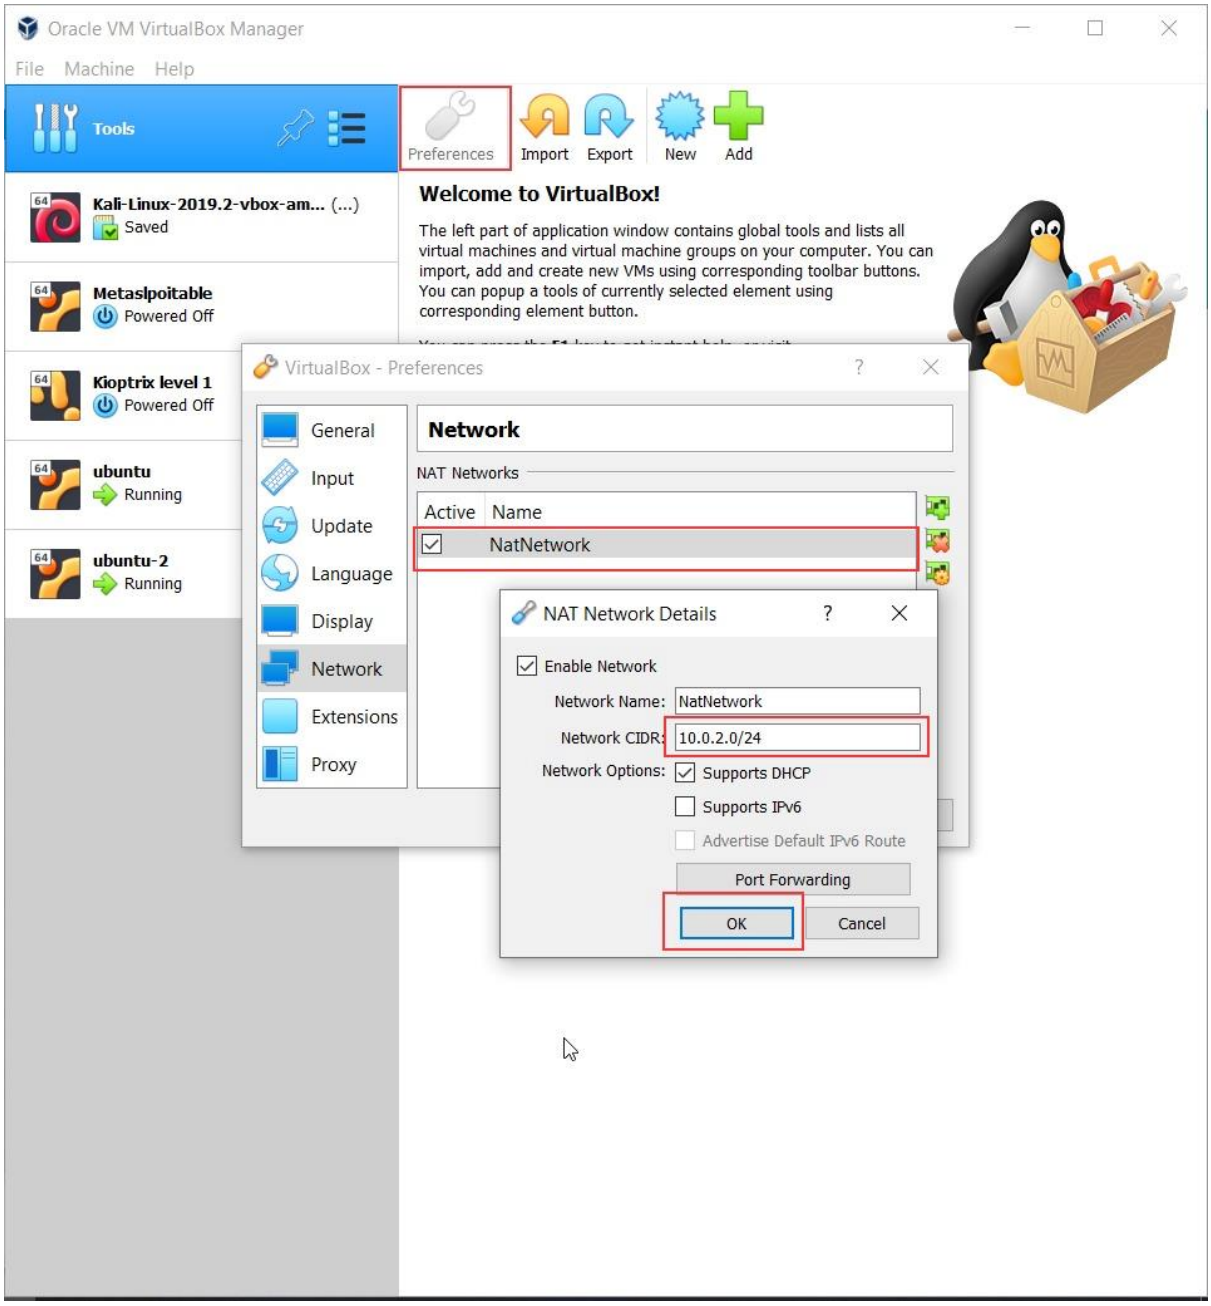

10.Find the IP address of the server machine

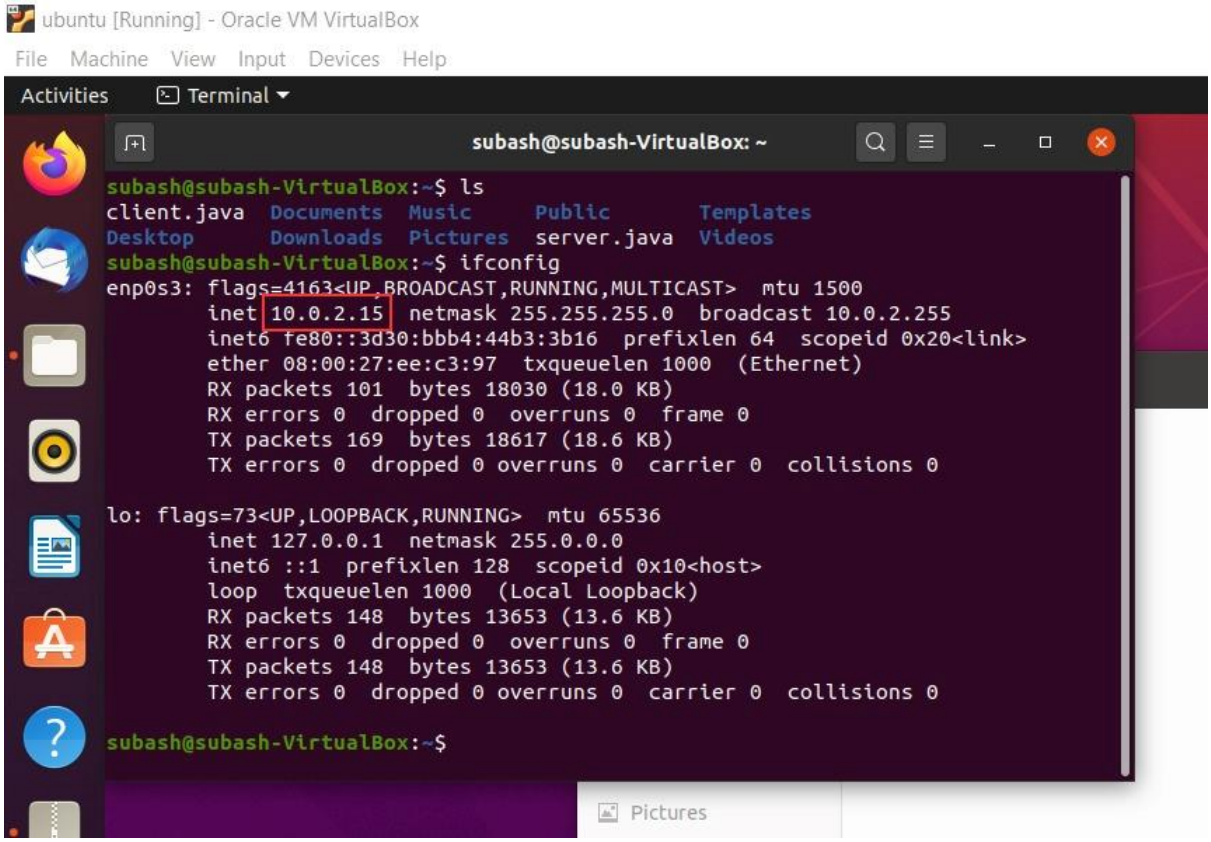

11.Use the IP address of the server machine in client program which is run in the client Machine

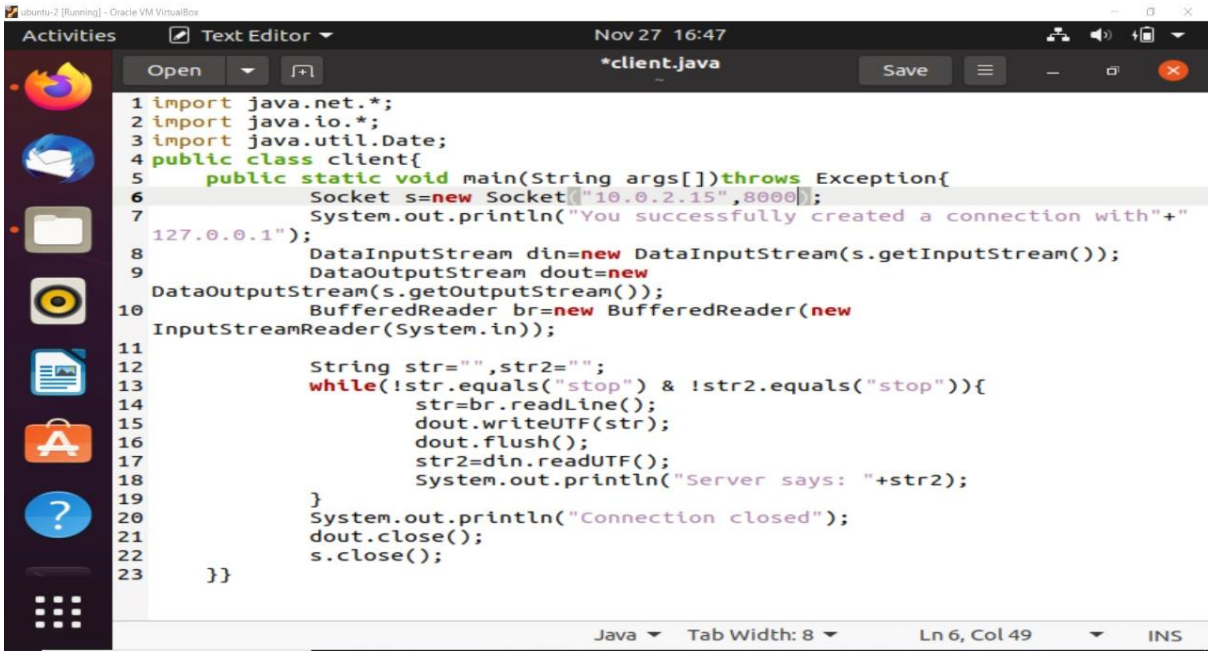

12.Configure the server program in the server machine by creating a server socket

Open server.java ~/.cache/.fr-4HBjPQ/exp6 Save

```
1 import java.net.*;
2 import java.io.*;
3 import java.util.Date;
4 public class server{
5     public static void main(String args[])throws Exception{
6             ServerSocket ss=new ServerSocket(8000);
7             Socket s=ss.accept();
8             System.out.println("client joined the conversation");
9             DataInputStream din=new DataInputStream(s.getInputStream());
10            DataOutputStream dout=new DataOutputStream(s.getOutputStream());
11            BufferedReader br=new BufferedReader(new InputStreamReader(System.in));
12            Date d=new Date();
13            String date="Current Date and time on sever is: " +d;
14            dout.writeUTF(date);
15           dout.flush();
16            String str="",str2="";
17            while(!str.equals("stop") & !str2.equals("stop")){
18                    str2=br.readLine();
19                    dout.writeUTF(str2);
20                    str=din.readUTF();
21                    System.out.println("vijay says: "+str);
22
23                    dout.flush();
24            }
25            System.out.println("Connection closed");
26            din.close();
27            s.close();
28            ss.close();
29    }}
```

Java Tab Width: 8 Ln 15, Col 25 INS

13.Compile and run the code in both machines for communication between both machines

```
ubuntu-2@ubuntu2-VirtualBox:~$ java client
You successfully created a connection with 127.0.0.1
hi
Server says: Current Date and time on sever is: Sat Nov 27 17:56:23 IST 2021
hi server
Server says: hi
i am vijay
Server says: hello vijay
```

```
^Csubash@subash-VirtualBox:~$ javac server.java
subash@subash-VirtualBox:~$ java server
Client joined the conversation
hi
vijay says: hi
hello vijay
vijay says: hi server
```

Experiment 2

Communication between the host machine and the virtual machine in the same computer

1.Configure the Nat Network

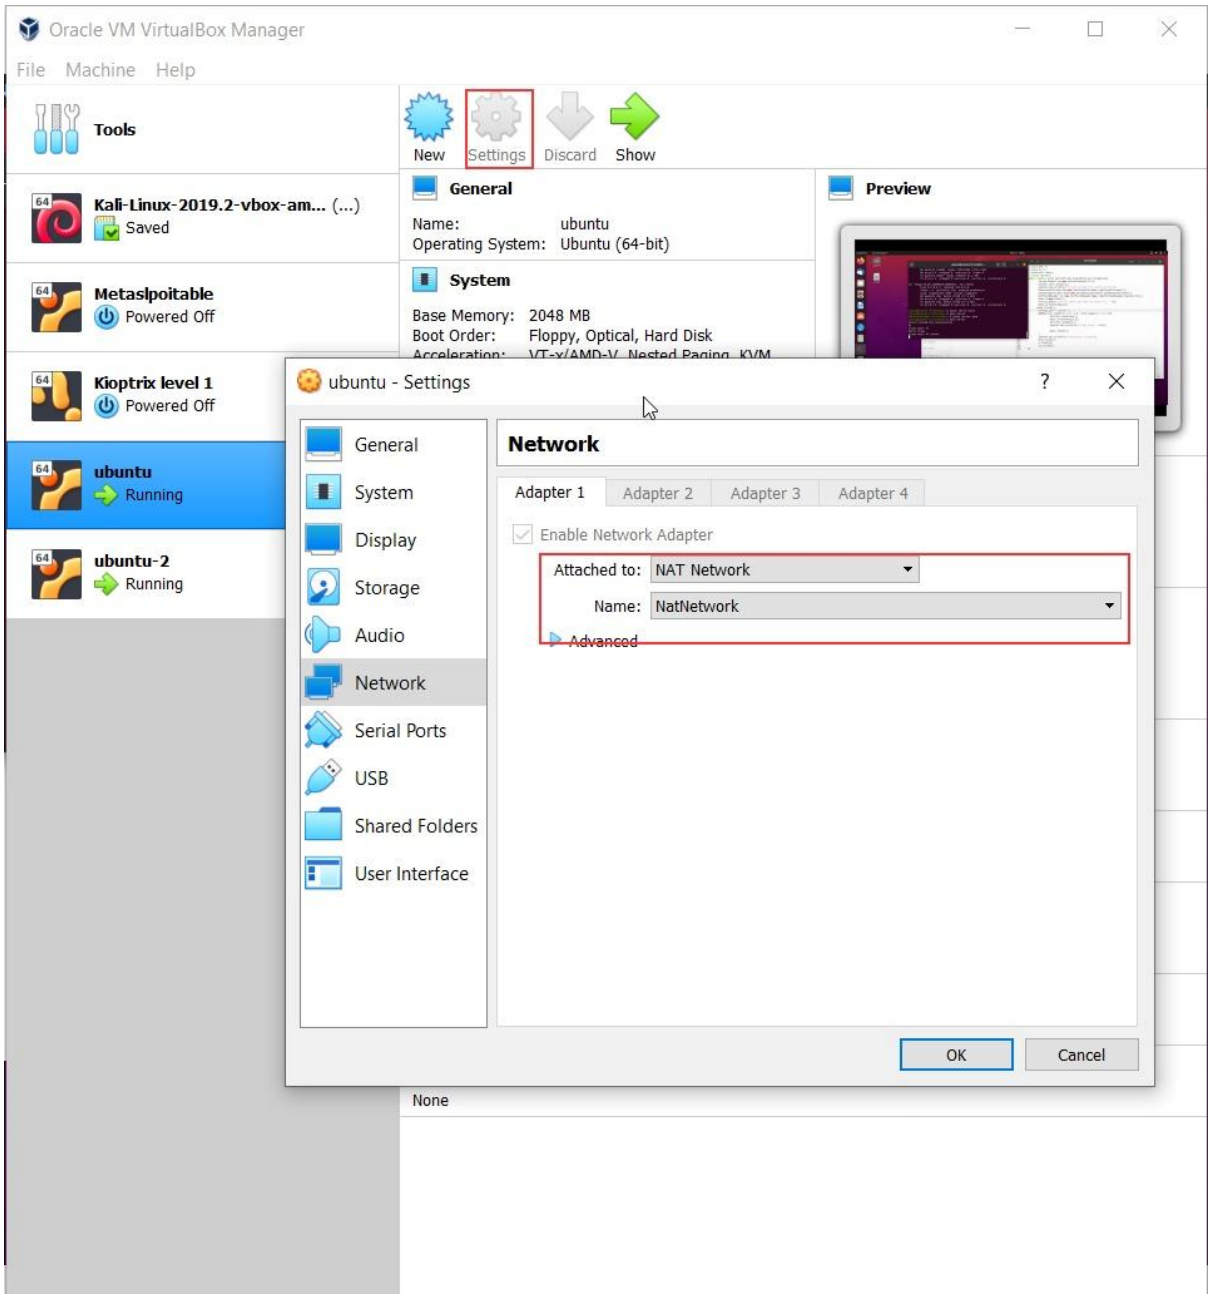

2.Do Port Forwarding to enable the communication between the machines

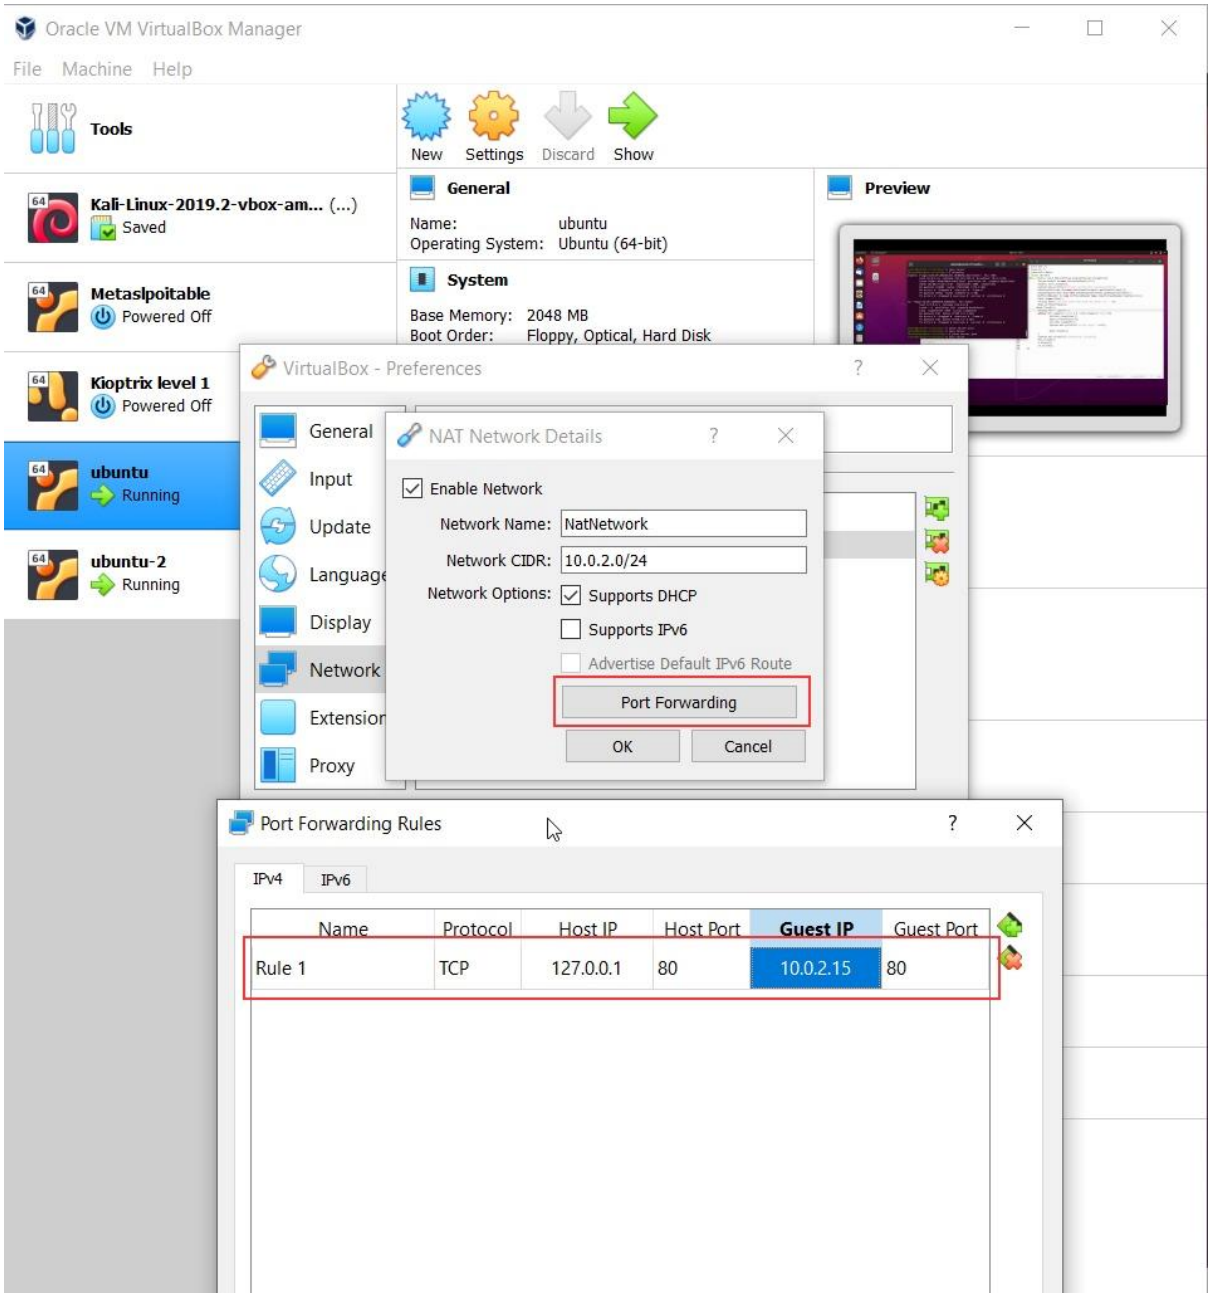

3.Create the server socket and run the server in the virtual machine

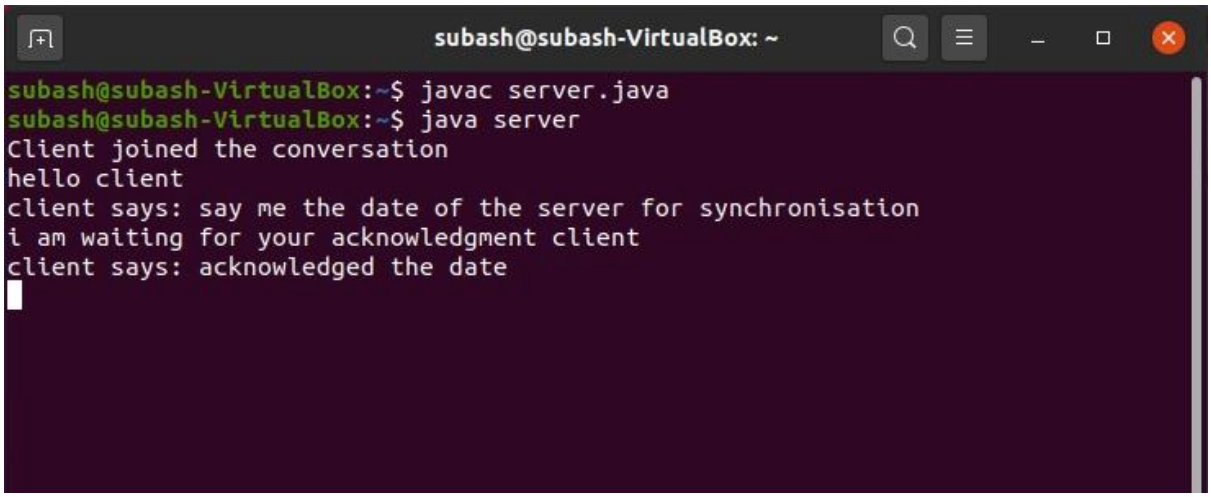

4.Open terminal in the host machine and connect the host with the virtual machine using the server machine's IP address and run the the client program.After the connection has been established you can perform the chat

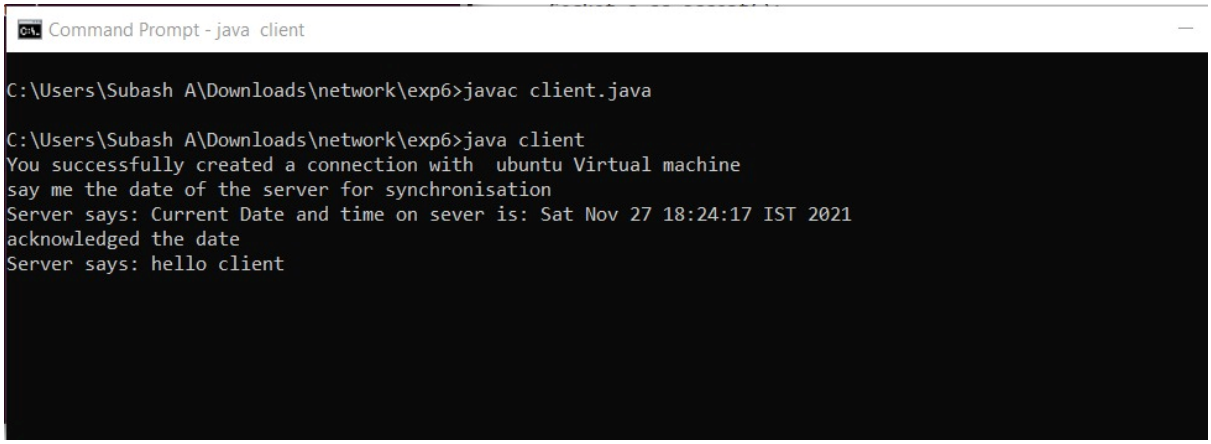

Experiment 3:

Communication between two virtual machines on different Hosts:

Proecdure:

1.Create two virtual machines on different hosts following the already mentioned steps

2.Comfigure the NAT Network and perform Port forwarding to connect both the machines

3.To perform this communication both the mamchines must be connected to the same network

4.Create Server Socket on one system and run the server program.

5.Run the client program on the other machine the servers port number and IP address.Thus connection would have established.Now you can perform the communication between both the machines

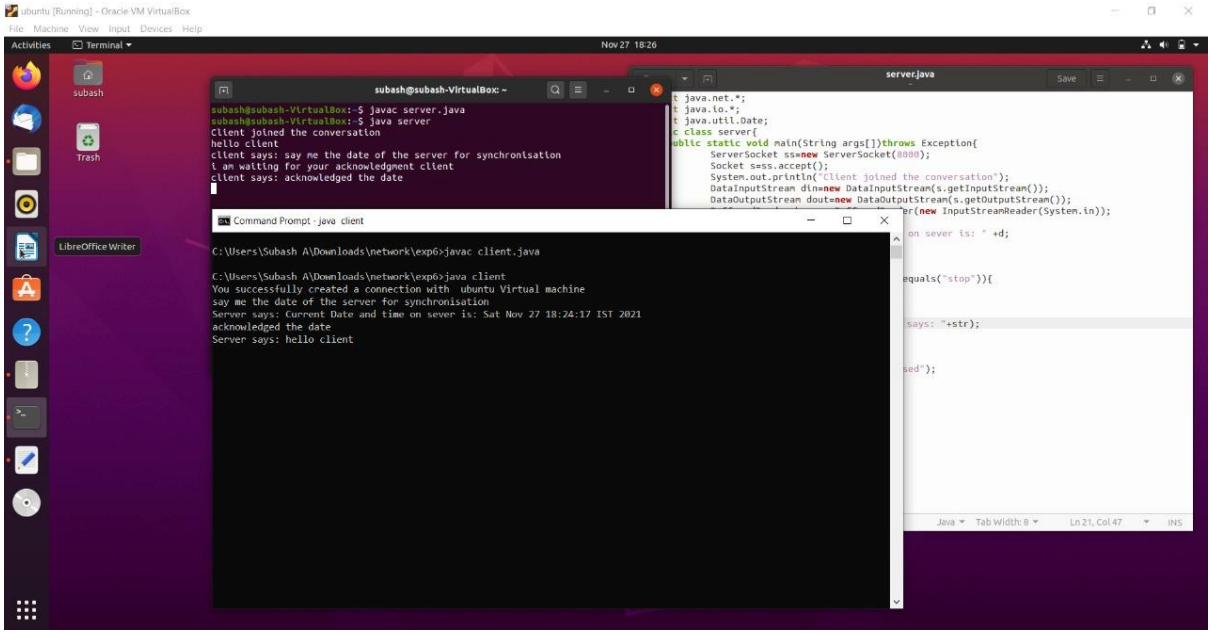
ubuntu [Running] - Oracle VM VirtualBox
File Machine View Input Devices Help
Activities
Terminal
Nov 27 18:26
subash
Trash
LibreOffice Writer
subash@subash-VirtualBox: ~
subash@subash-VirtualBox:~$ javac server.java
subash@subash-VirtualBox:~$ java server
Client joined the conversation
hello client
client says: say me the date of the server for synchronisation
i am waiting for your acknowledgment client
client says: acknowledged the date
Command Prompt - java client
C:\Users\Subash A\Downloads\network\exp6>javac client.java
C:\Users\Subash A\Downloads\network\exp6>java client
You successfully created a connection with ubuntu Virtual machine
say me the date of the server for synchronisation
Server says: Current Date and time on sever is: Sat Nov 27 18:24:17 IST 2021
acknowledged the date
Server says: hello client
server.java
Save
Java
Tab Width: 8
Ln 21, Col 47
INS

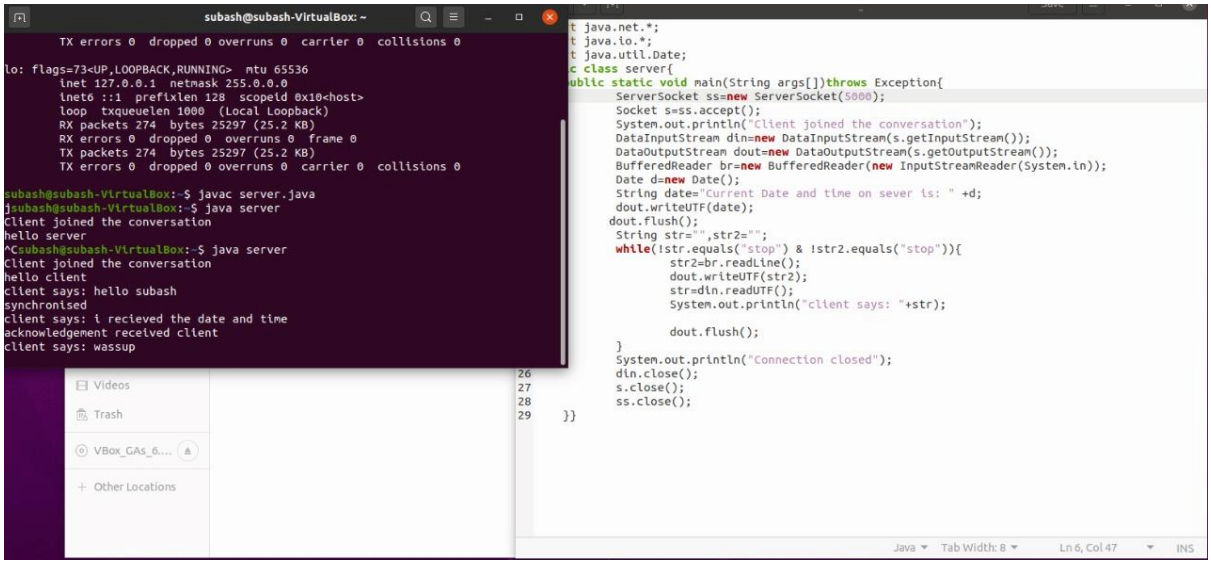
subash@subash-VirtualBox: ~
TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
inet 127.0.0.1 netmask 255.0.0.0
inet6 ::1 prefixlen 128 scopeid 0x10<host>
loop txqueuelen 1000 (Local Loopback)
RX packets 274 bytes 25297 (25.2 KB)
RX errors 0 dropped 0 overruns 0 frame 0
TX packets 274 bytes 25297 (25.2 KB)
TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
subash@subash-VirtualBox:~$ javac server.java
subash@subash-VirtualBox:~$ java server
Client joined the conversation
hello server
^Csubash@subash-VirtualBox:~$ java server
Client joined the conversation
hello client
client says: hello subash
synchronised
client says: i recieved the date and time
acknowledgement received client
client says: wassup
Videos
Trash
VBox_GAs_6....
Other Locations
t java.net.*;
t java.io.*;
t java.util.Date;
c class server{
ublic static void main(String args[])throws Exception{
ServerSocket ss=new ServerSocket(5000);
Socket s=ss.accept();
System.out.println("Client joined the conversation");
DataInputStream din=new DataInputStream(s.getInputStream());
DataOutputStream dout=new DataOutputStream(s.getOutputStream());
BufferedReader br=new BufferedReader(new InputStreamReader(System.in));
Date d=new Date();
String date="Current Date and time on sever is: " +d;
dout.writeUTF(date);
dout.flush();
String str="",str2="";
while(!str.equals("stop") & !str2.equals("stop")){
str2=br.readLine();
dout.writeUTF(str2);
str=din.readUTF();
System.out.println("client says: "+str);
dout.flush();
}
System.out.println("Connection closed");
din.close();
s.close();
ss.close();
}}
Java
Tab Width: 8
Ln 6, Col 47
INS

Experiment4.

Map reduce:

Procedure:

1.Create new User.

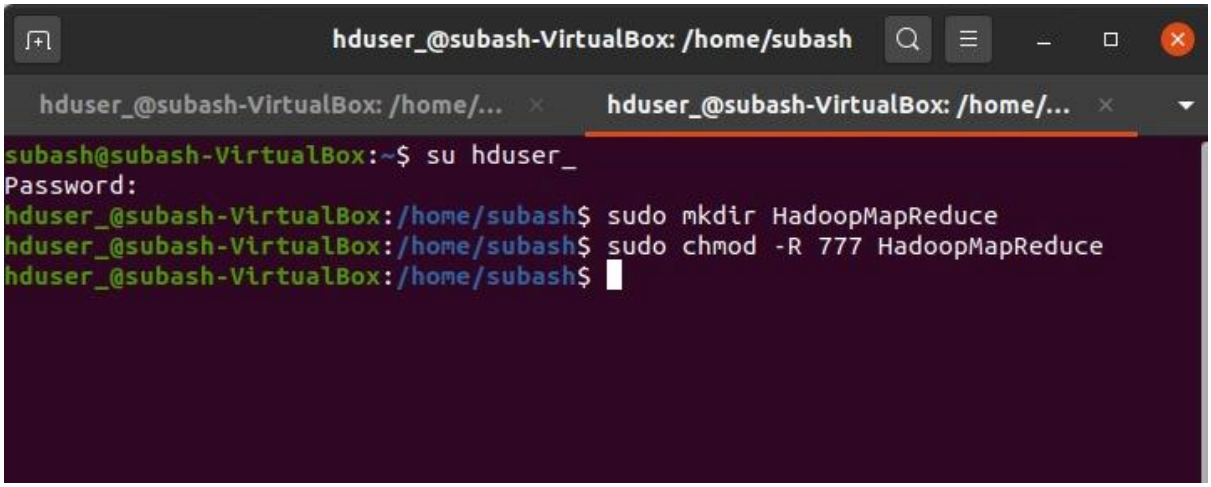

2.We write the programs for the mapping,shuffling and reducing.The programs are shown below.

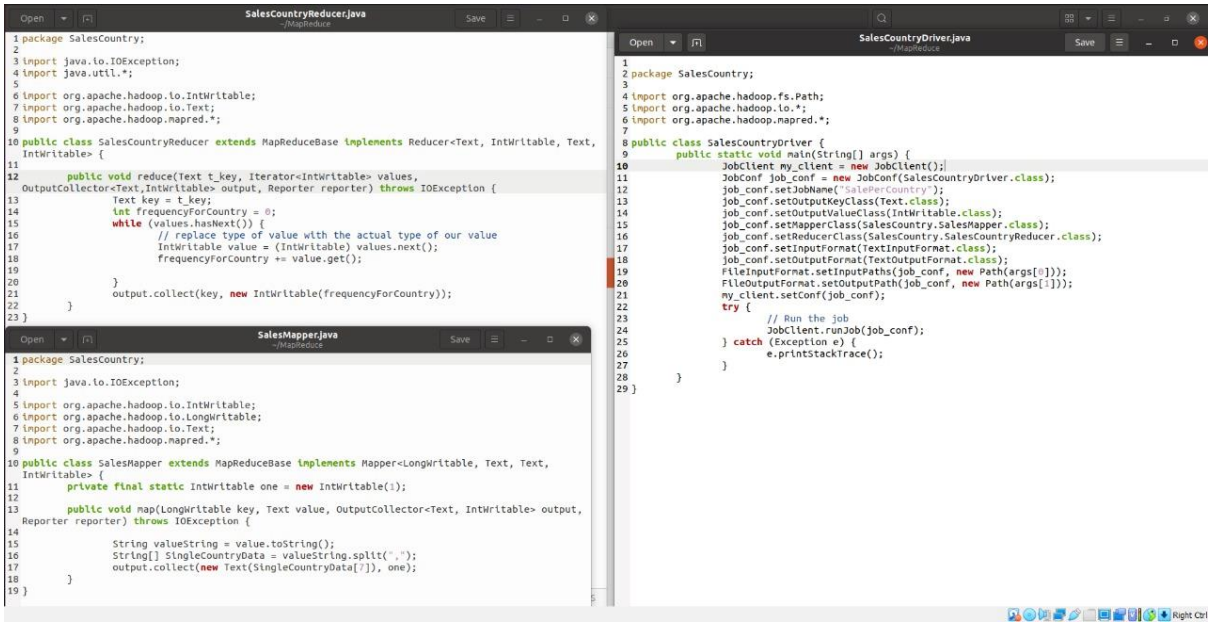

3.Below is the csv File which needs to be mapreduced.

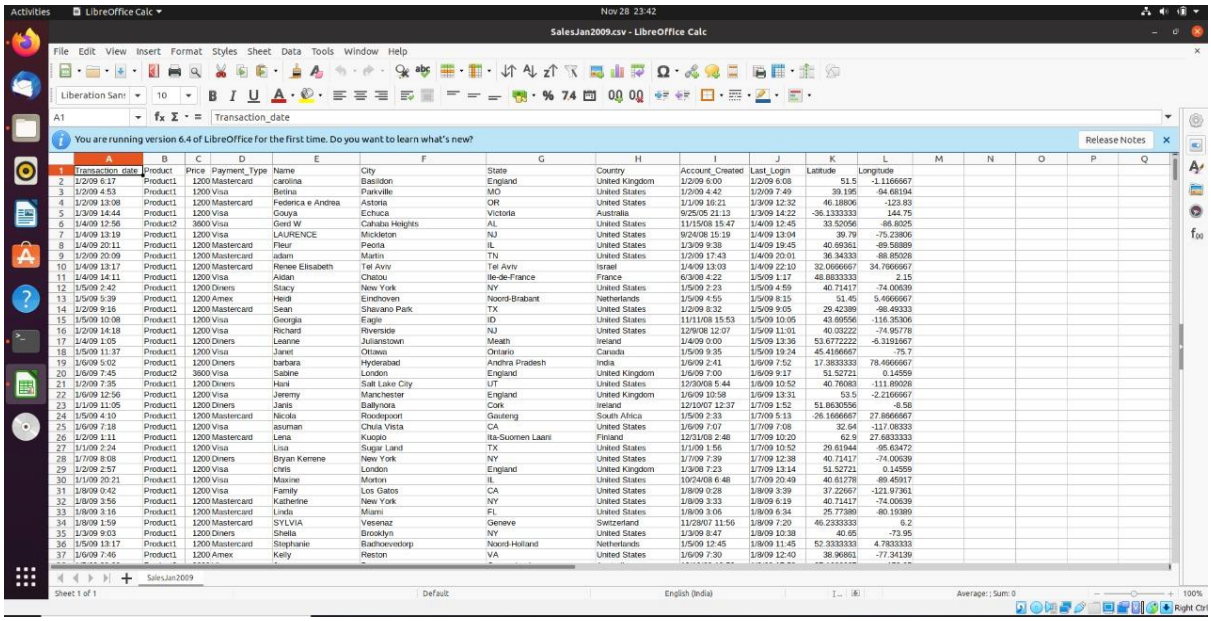

4.In the command Prompt we gert Hadoop started.

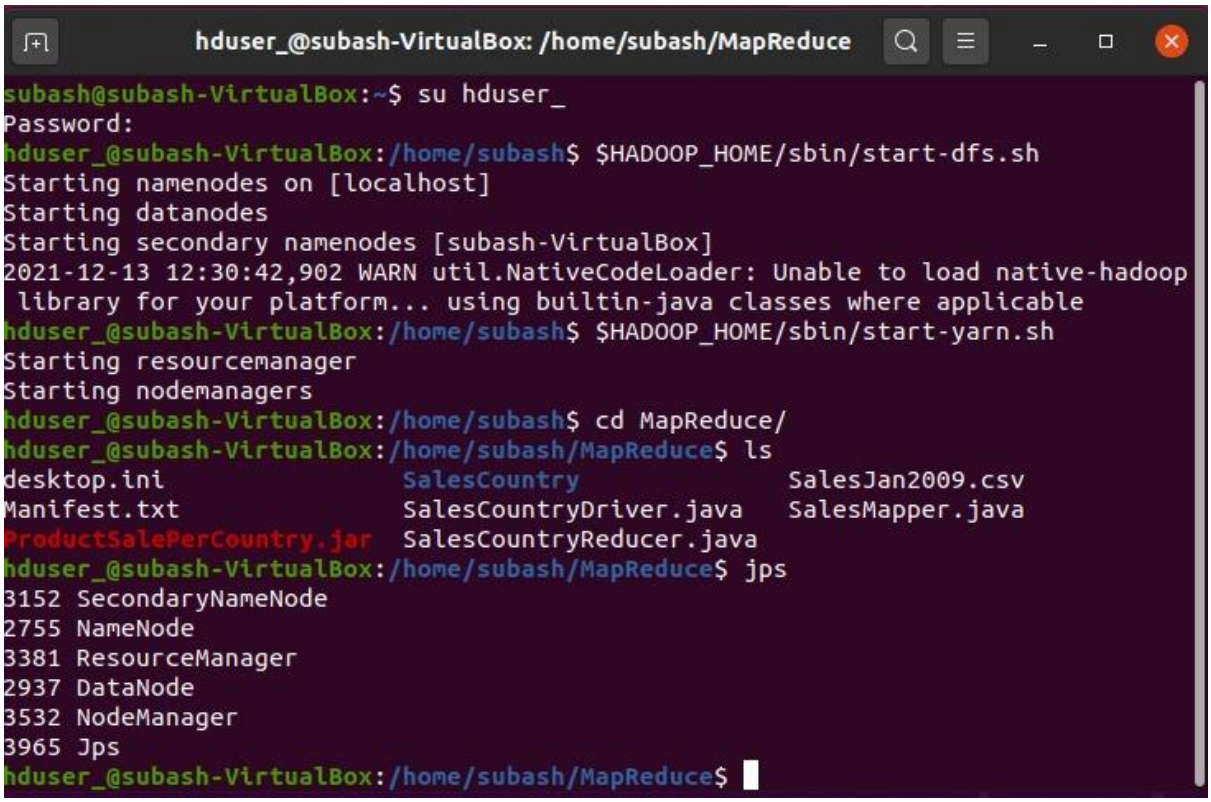

5.Then we run the jar file which has the compilation of the programs required for map reduce.

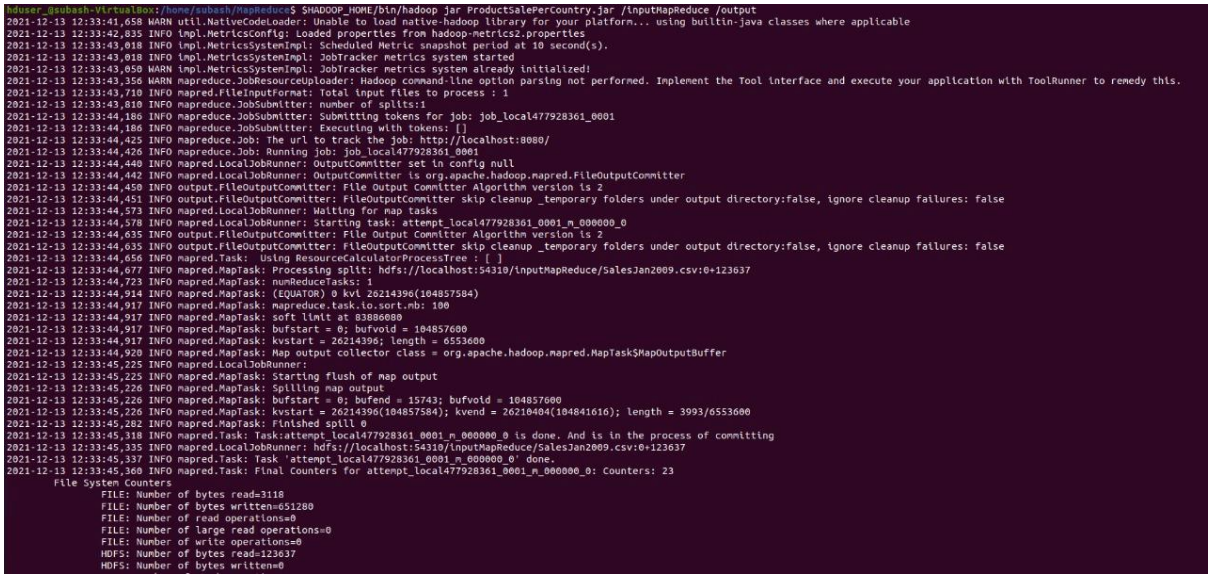

6. We can see the process happening

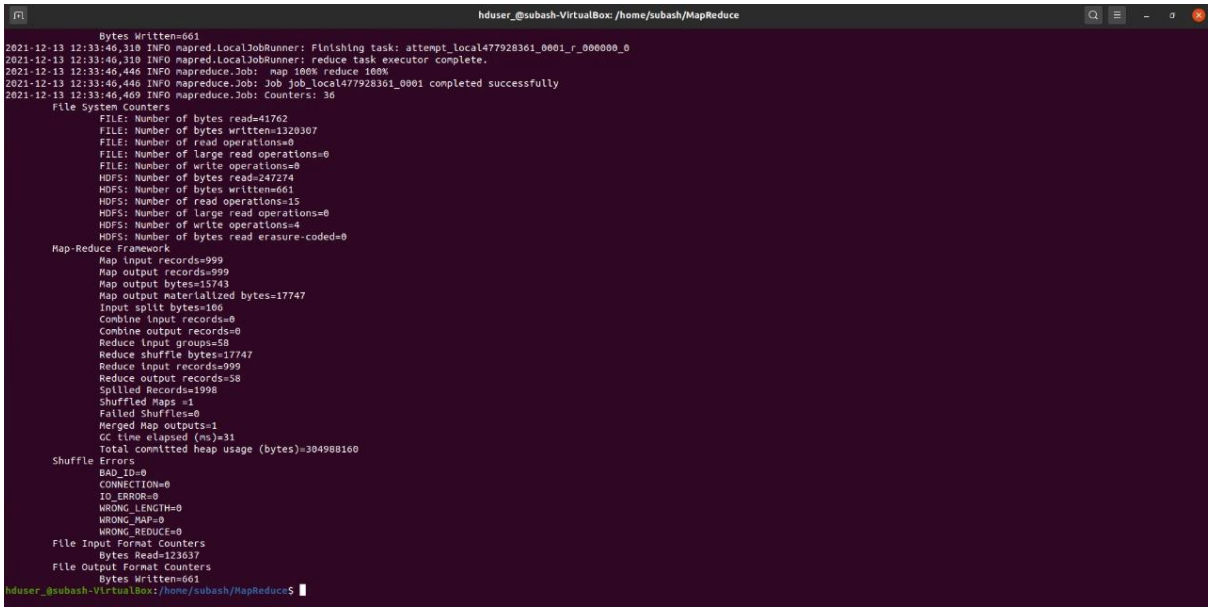

7.From the directory we specified for the output,we could access the final output.We are left with the reduced data.

Concurrent Programming

About concurrency:

Concurrency is the ability of a processor to execute multiple instruction sequences to run on a single/multiple processor, thus using the full processing capability of a processor. It allows for parallel execution of concurrent units which can significantly improve the overall speed of execution in multi-processor and multi-core system. Although concurrency allows parallel execution, it is slightly different from parallelism. Concurrency is about dealing with lot of things at once whereas parallelism is about doing a lot of things at once. Basically concurrency is about and structure whereas parallelism is about execution.

Concurrency is achieved in Java by using Threads and the Runnable interface. Multiple threads are spawned to execute the multiple instruction sequences and each thread's execution is controlled by the Runnable interface. Synchronization is achieved in Java using the synchronized keyword which allows the function(shared resource) to be executable only by one thread at a time.

Multiple clients – server chat in Java using concurrency:

Server is used here to wrap each client with a thread using a ClientHandler class and show how many clients are connected at that time of instance. ClientHandler class is used to get the incoming messages from all clients and display it on the current client's terminal. The Client class is used to get inputs from each client terminal and send it to the ClientHandler class which would broadcast the message to other clients. Each client is identified by their username so that we can tell the other clients who is sending the particular message.

The code goes as follows:

Server.java

```
import java.io.IOException;
import java.net.ServerSocket;
import java.net.Socket;

public class Server {
    private ServerSocket ss;

    public Server(ServerSocket ss) {
        this.ss = ss;
    }

    public void startServer() {
        int count = 0;
        try {
            while(!ss.isClosed()) {
                Socket s = ss.accept();
                System.out.println(++count +" client(s) connected");
                ClientHandler clientHandler = new ClientHandler(s); // ClientHandler implements Runnable which will spawn a new thread for each client that is connected with the server

                Thread thread = new Thread(clientHandler); // Thread created to invoke the run() in ClientHandler
                thread.start();
            }
        } catch (IOException e) {
```

```
            closeServerSocket();
        }
    }

    public void closeServerSocket() { // function to avoid nested try catches in
startServer
        try {
            if(ss != null) {
                ss.close();
            }
        } catch (IOException e) {
            e.printStackTrace();
        }
    }

    public static void main(String[] args) throws IOException {
        ServerSocket serverSocket = new ServerSocket(6767);
        Server server = new Server(serverSocket);
        server.startServer();
    }
}
```

Client.java

```
import java.io.*;
import java.net.Socket;
import java.util.Scanner;
```

```
public class Client {
    private Socket s;
    private BufferedReader in;
    private BufferedWriter out;
    private String clientName;

    public Client(Socket s, String clientName) {
        try {
            this.s = s;
            this.in = new BufferedReader(new
InputStreamReader(s.getInputStream()));
            this.out = new BufferedWriter(new
OutputStreamWriter(s.getOutputStream()));
            this.clientName = clientName;
        } catch(IOException e) {
            closeSocketAndStreams(s, in, out);
        }
    }

    public void sendMessage() {
        try {
            out.write(clientName);
            out.newLine();
            out.flush();

            Scanner scanner = new Scanner(System.in);
            while(s.isConnected()) {
                String messageToSend = scanner.nextLine();
                out.write(clientName+": "+messageToSend);
```

```
                out.newLine();
                out.flush();
            }
        } catch (IOException e) {
            closeSocketAndStreams(s, in, out);
        }
    }

    public void listenForMessages() {
        new Thread(new Runnable() { // listening to messages is a blocking operation. If thread is not used here then we may end up waiting for messages from other users and not be able to send messages
            @Override
            public void run() {
                String messageFromClients;

                while(s.isConnected()) {
                    try {
                        messageFromClients = in.readLine();
                        System.out.println(messageFromClients);
                    } catch(IOException e) {
                        closeSocketAndStreams(s, in, out);
                    }
                }
            }
        }).start();
    }
```

```
    public void closeSocketAndStreams(Socket s, BufferedReader in,
BufferedWriter out) {

        try {
            if(in != null) {
                in.close();
            }

            if(out != null) {
                out.close();
            }

            if(s != null) {
                s.close();
            }
        } catch(IOException e) {
            e.printStackTrace();
        }
    }

    public static void main(String[] args) throws IOException {
        Scanner scanner = new Scanner(System.in);
        System.out.println("Enter your name");
        String clientName = scanner.nextLine();
        Socket s = new Socket("localhost", 6767);
        Client client = new Client(s, clientName);
        client.listenForMessages();
        client.sendMessage();
    }
```

```
}
```

ClientHandler.java

```
import java.io.*;
import java.net.Socket;
import java.util.ArrayList;

public class ClientHandler implements Runnable{

    public static ArrayList<ClientHandler> clientHandlers = new ArrayList<>();
    private Socket s;
    private BufferedReader in;
    private BufferedWriter out;
    private String clientName;

    public ClientHandler(Socket s) {
        try {
            this.s = s;
            this.out = new BufferedWriter(new
OutputStreamWriter(s.getOutputStream()));
            this.in = new BufferedReader(new
InputStreamReader(s.getInputStream()));
            this.clientName = in.readLine();//get the client's name using
BufferedReader
            clientHandlers.add(this);
            broadcastMessage("Server: "+clientName+" has joined");
        } catch (IOException e) {
```

```
            closeSocketAndStreams(s, in, out);
        }
    }


    @Override
    public void run() {
        String messageFromClient;
        while(s.isConnected()) {
            try {
                messageFromClient = in.readLine(); // reading a message will block the client thread. Therefore it must be run on seperate threads
                if(messageFromClient.equals("stop")) {
                    throw new IOException(clientName+" has left");
                } else {
                    broadcastMessage(messageFromClient);
                }

            } catch (IOException e) {
                closeSocketAndStreams(s, in, out);
                break;
            }
        }
    }

    public void broadcastMessage(String messageToSend) {
        for(ClientHandler clientHandler: clientHandlers) {
            try {
                if(!clientHandler.clientName.equals(clientName)) {
```

```
                clientHandler.out.write(messageToSend);
                clientHandler.out.newLine();
                clientHandler.out.flush();
            }
        } catch(IOException e) {
            closeSocketAndStreams(s, in, out);
        }
    }
}

public void removeClientHandler() {
    clientHandlers.remove(this);
    broadcastMessage("Server: "+clientName+" has left");
}

public void closeSocketAndStreams(Socket s, BufferedReader in,
BufferedWriter out) {
    removeClientHandler(); // if some error happens with a client we are
removing the client from the chat
    try {
        if(in != null) {
            in.close();
        }

        if(out != null) {
            out.close();
        }

        if(s != null) {
```

```
                s.close();
            }
        } catch(IOException e) {
            e.printStackTrace();
        }
    }
}
```

Output:

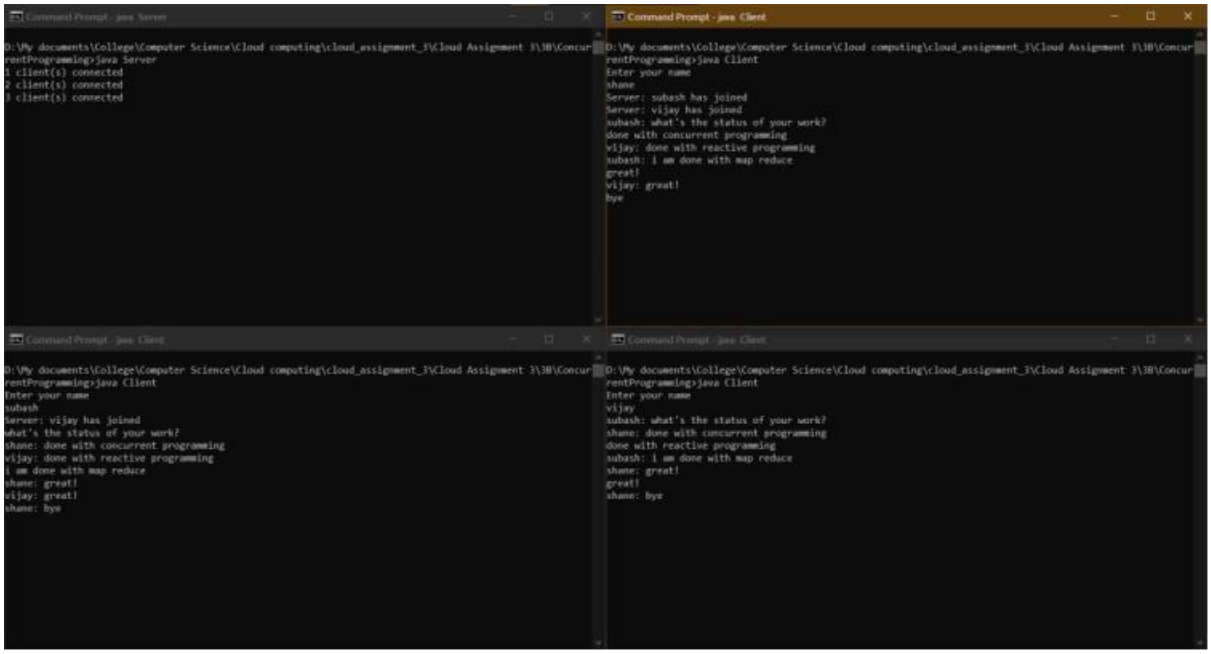

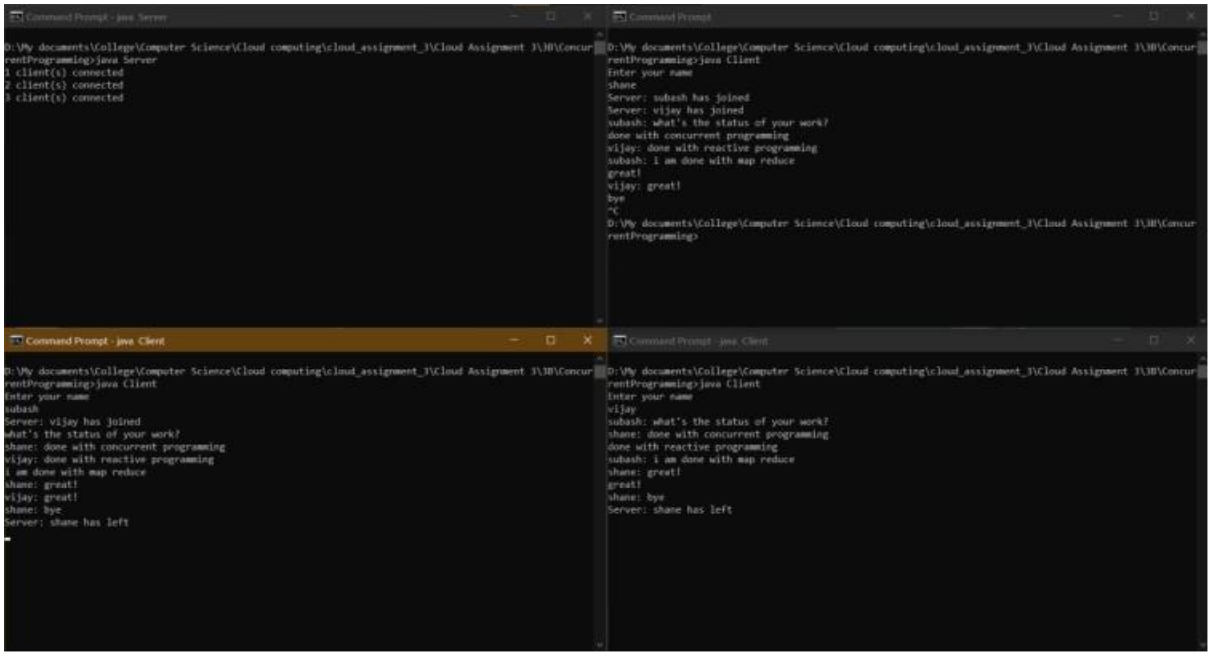

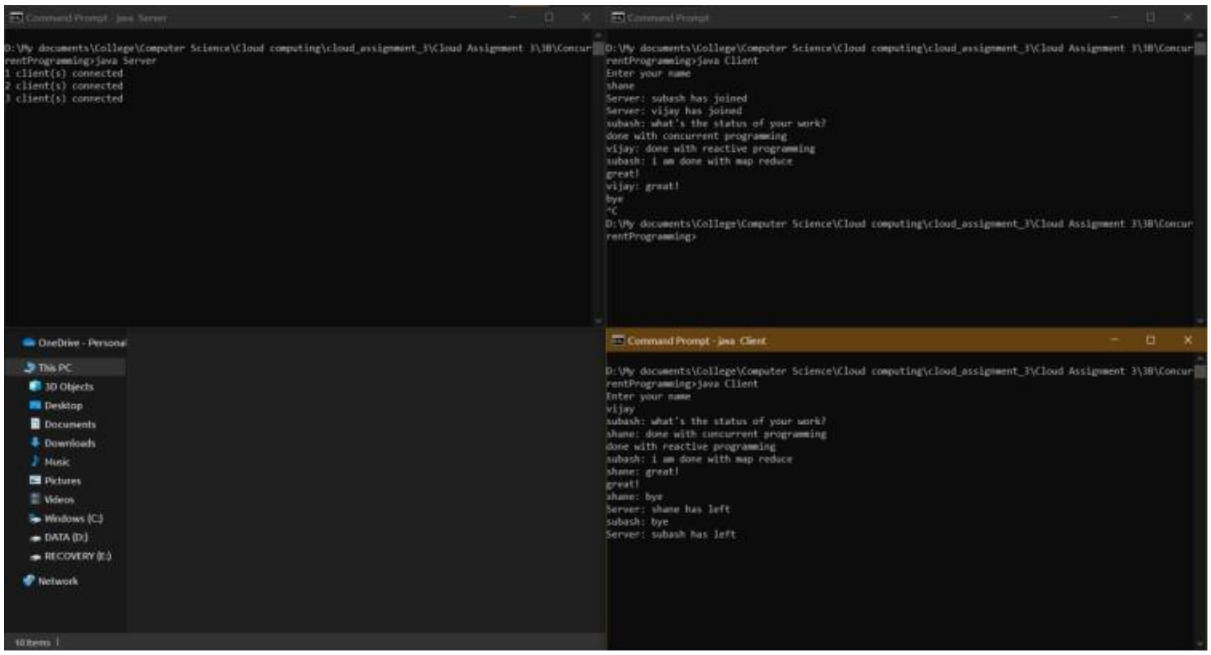

Thus, concurrent programming model has been used to perform multiple clients server chat successfully.

Reactive Programming using RxJS

Step 1: To work with RxJS, we need the following things: NodeJS, npm and RxJS package installation

Step 2: Download NodeJS from the website https://nodejs.org/en/download/. Installation of NodeJS is similar to other free and open source software. This would install npm as well

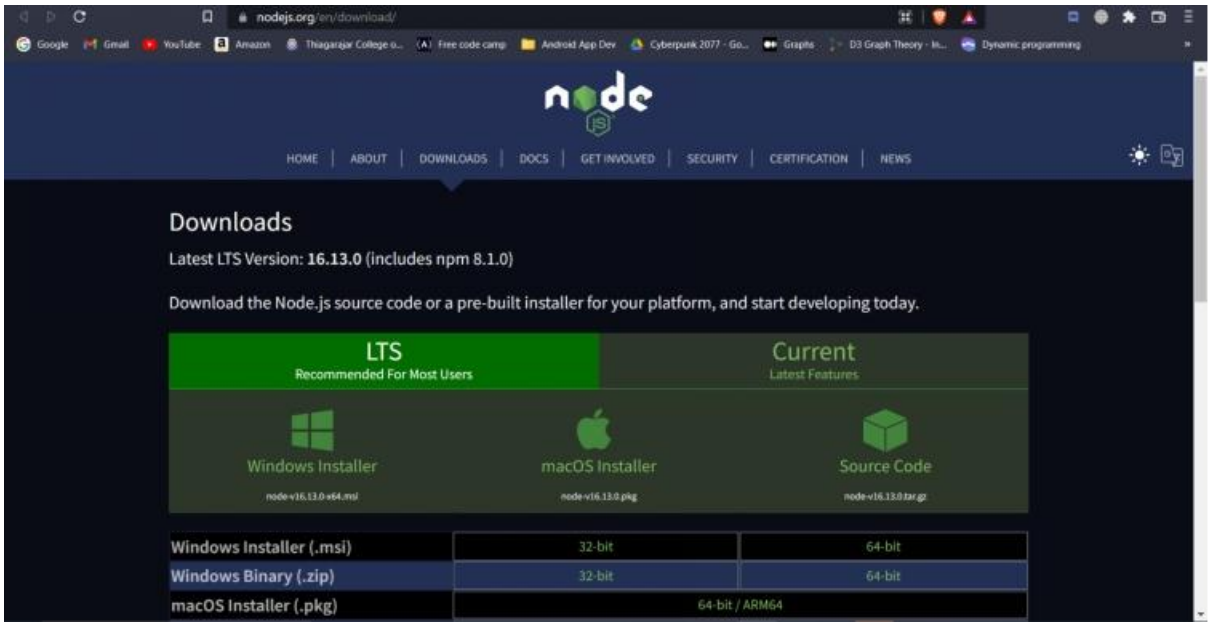

Step 3: Create a folder named rxjsproj to store the example program and navigate to that directory in cmd. Then type npm init to create a package.json for project setup

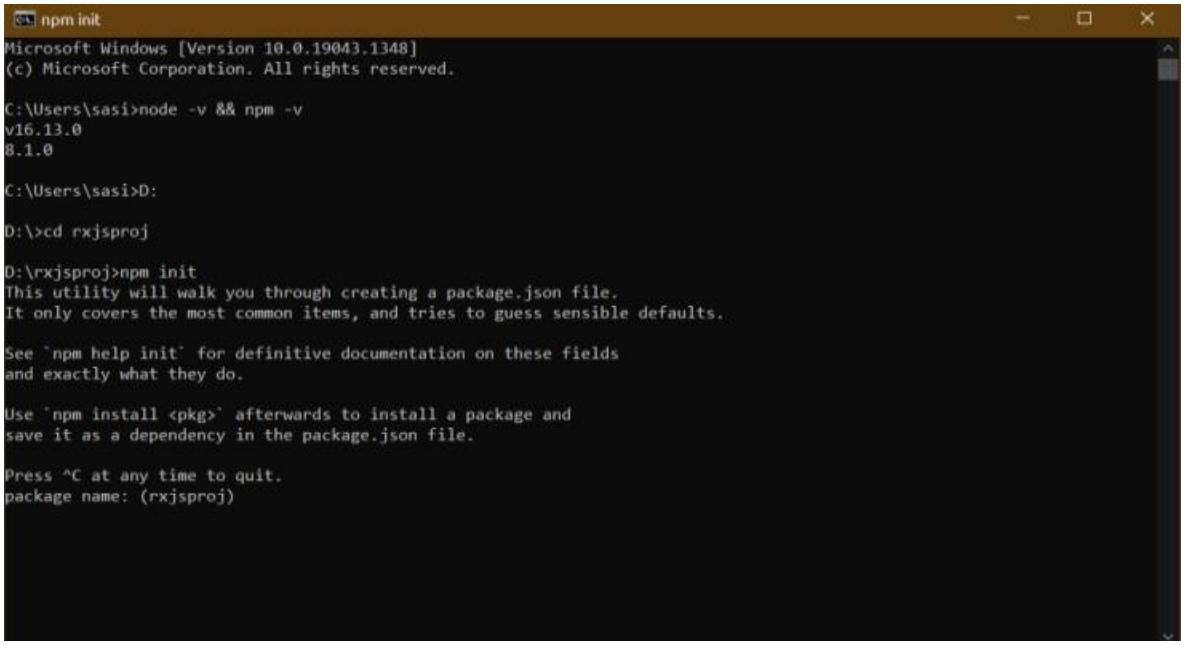

Step 4: After this process you will see a package.json file in the rxjsproj folder

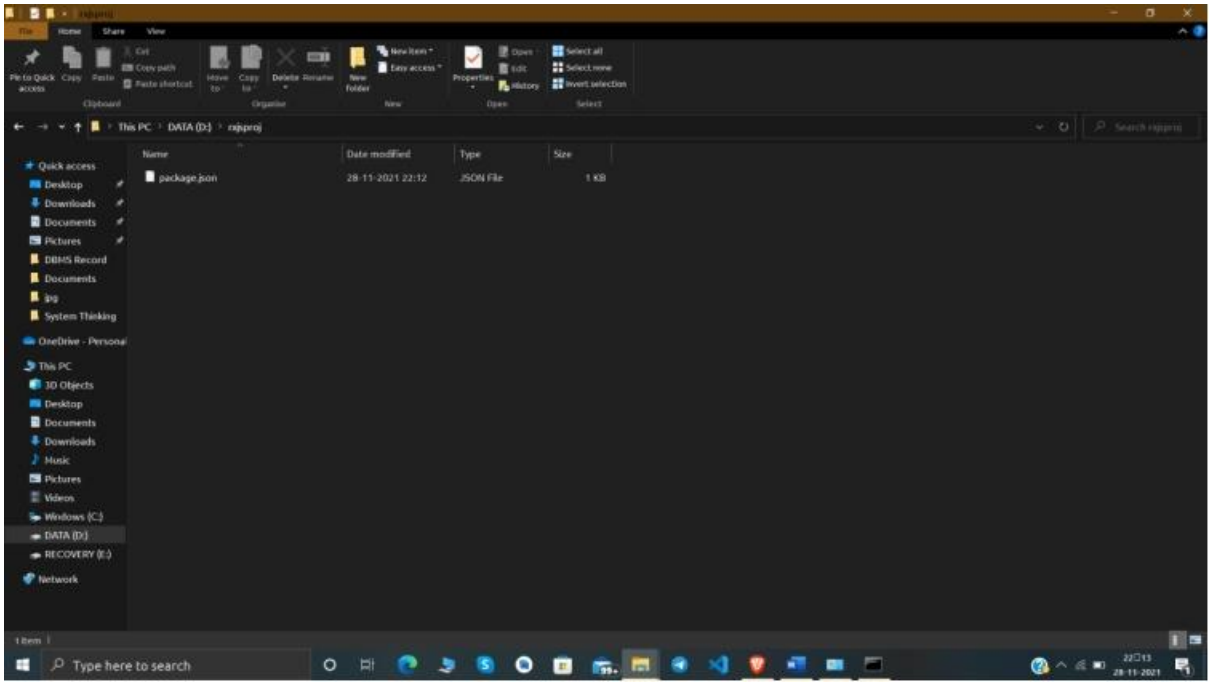

Step 5: Install rxjs using the following command: npm install ---save-dev rxjs

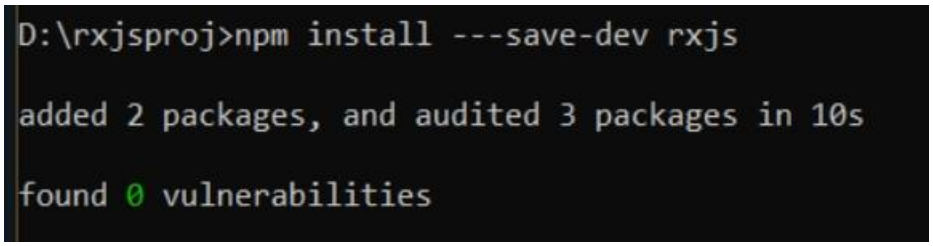

Step 6: Install ep6 modules to avoid any problems with importing packages which node requires using the following command: npm install --save-dev esm

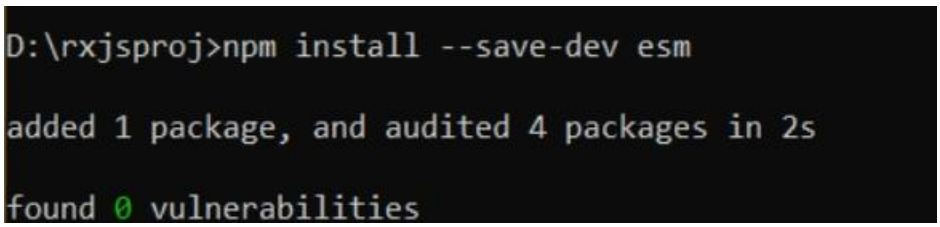

Step 7: Create a javascript file named square.js to demonstrate reactive programming

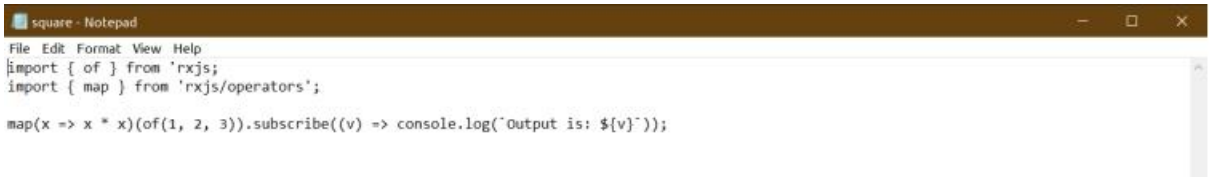

Step 8: Run the program using the following command: node -r esm square.js

And we get the desired output after execution